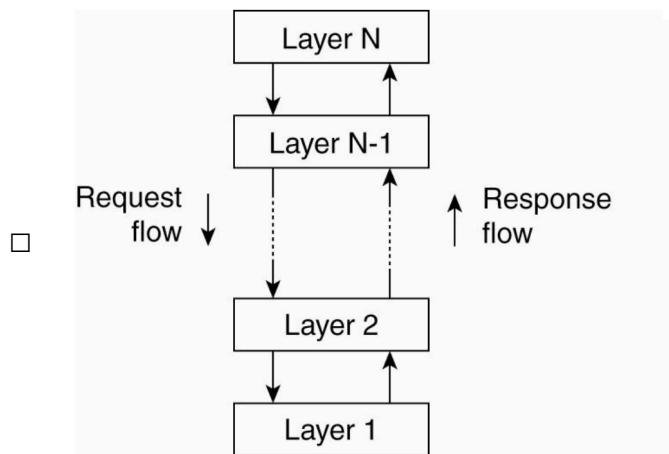


# Lecture 1

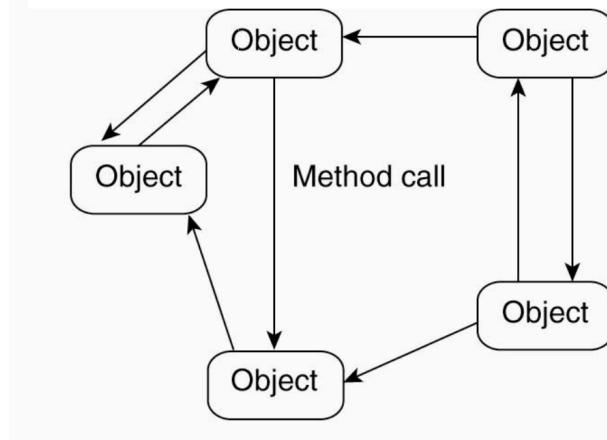
Monday, 7 August 2017 2:28 PM

## Distributed system and network principle

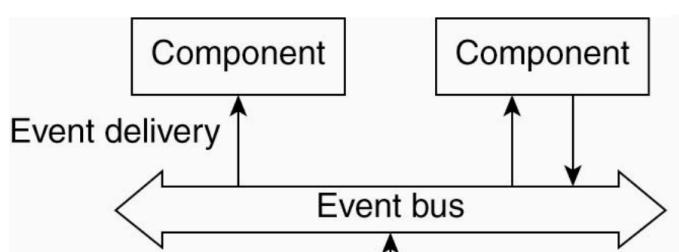
- Software architecture:
  - o Component organization
    - Layer-based architecture requests (resp. response) do downward (resp. upward) flow between layers
    - Network-based architecture communicate through Remote Procedure Call (RPC)



- Network-based architecture communicate through Remote Procedure Call (RPC)



- o Communication organization:
  - Event-based architecture: communication through events, that optionally involve an event bus

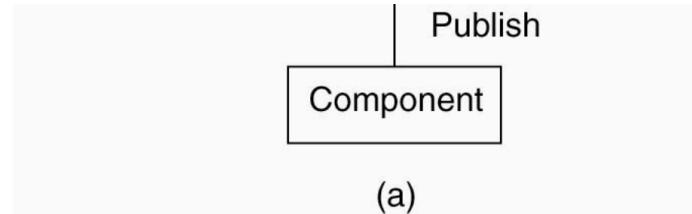


pward)

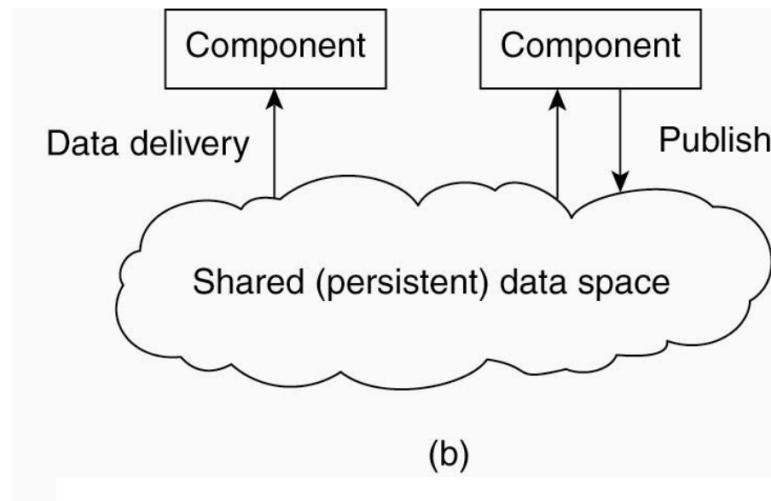
lls (RPCs)

carry data (subscribers get their desired events delivered)





- Data-centered architecture: through a shared repository, that contains data



- The client-server model:
  - The client server model:
    - The basic client -server model:
      - The client requests the service whereas the server provides the service
      - The client and the server can be hosted on different machines
    - Stateless and stateful server:
      - Stateless server: does not record the state of its clients
      - Stateful server: maintains persistent information about its clients (c)

	<b>Stateless server</b>	<b>Stateful server</b>
<b>State</b>	No info kept	Persistent info
<b>Request</b>	Self-contained	Can be split, generally faster
<b>Upon failure</b>	No recovery needed	State recovery needed (explicit)
<b>Example</b>	Network file system (NFSv3)	Andrew file system (AFS)

- The layered organization
  - Application layering:
    - Traditional three-layered views:
      - The user interface layer: contains the feature to control the application
      - The processing layer: contains the function of the application
      - The data layer: contains the data of the application
    - Example: a search engine request spanning the traditional three layers

ta (e.g. files in a distributed file system)

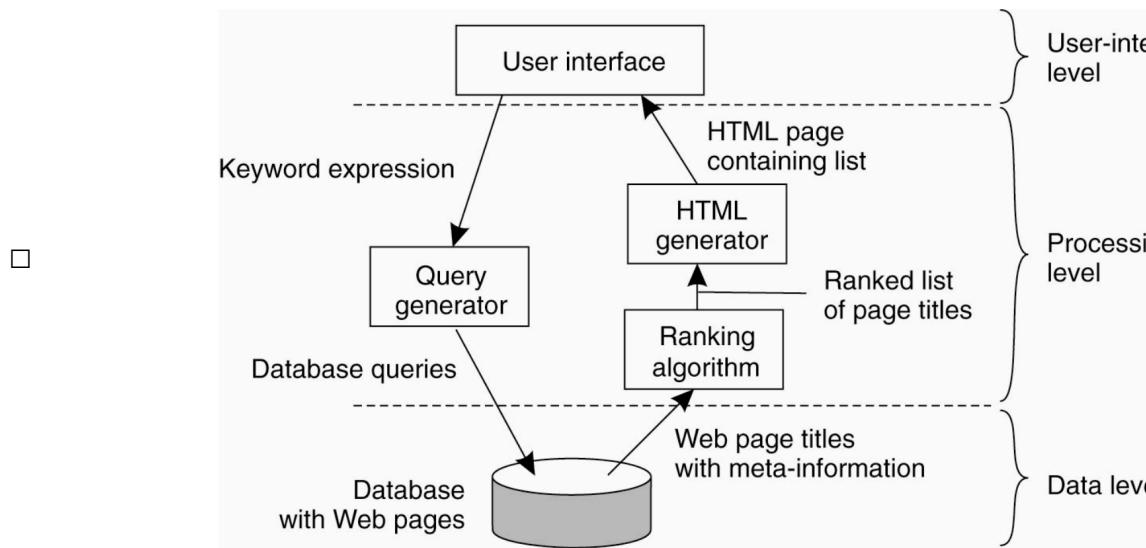
ice

lient -->file)

deletion)

tion

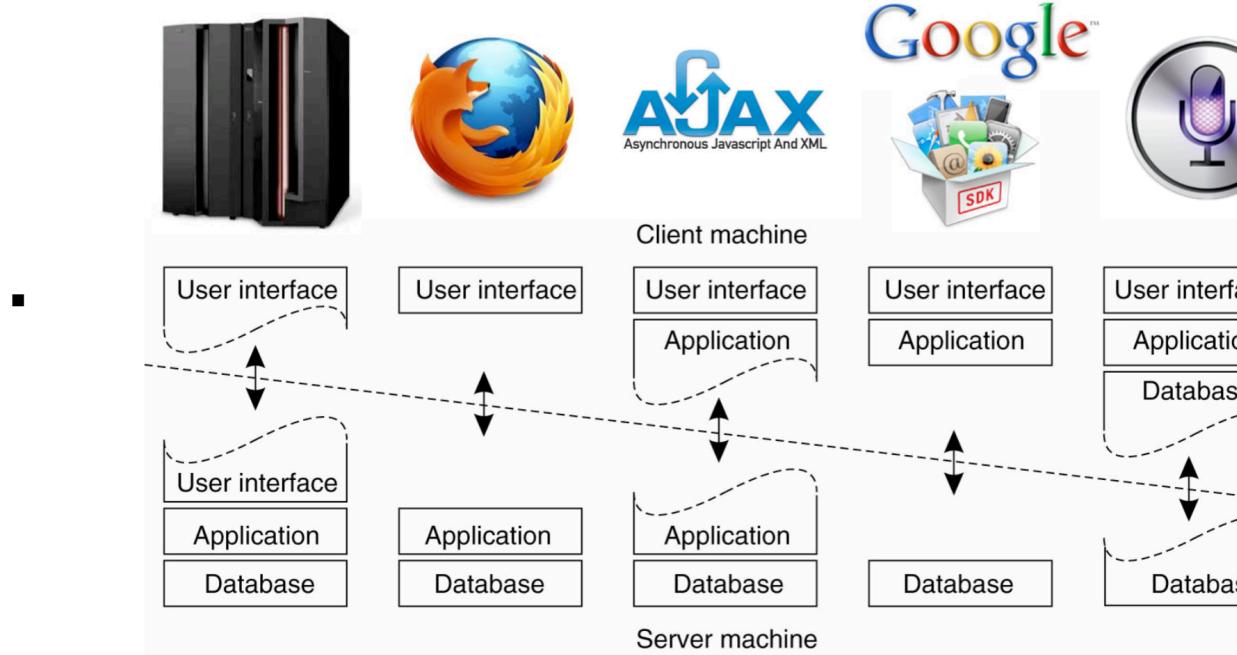




- Hosting different layers on different machines:
  - Three-tiered architecture:
    - ◆ each layer on separate machine
  - Two-tiered architecture:
    - ◆ Client
    - ◆ Single server configuration
  - Single-tiered architecture:
    - ◆ Dumb terminal
    - ◆ Mainframe configuration

- Multi-tiered architectures:

### Physical two-tiered architecture



A single machine can act both as a client and a server

erface

ng

el

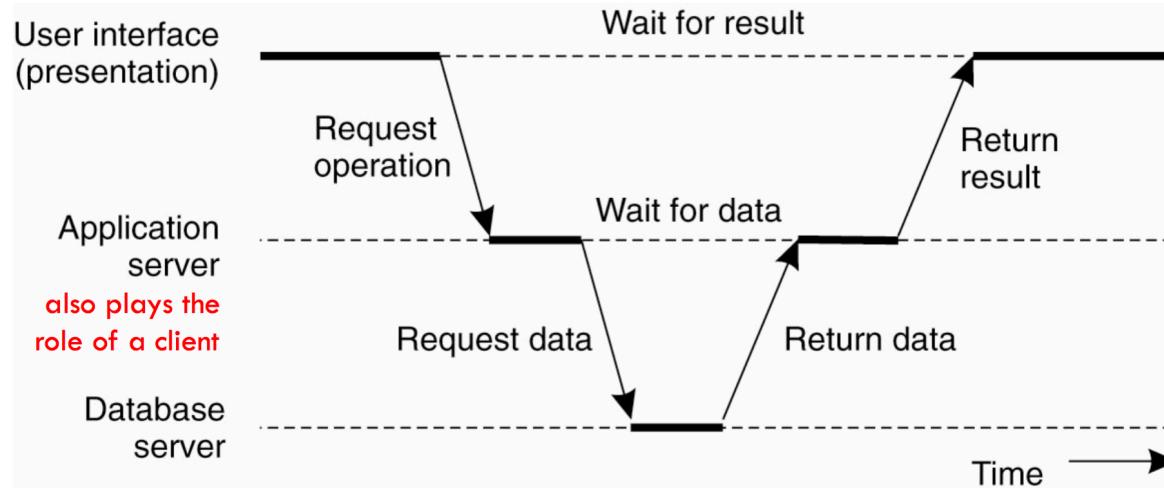


ace  
on

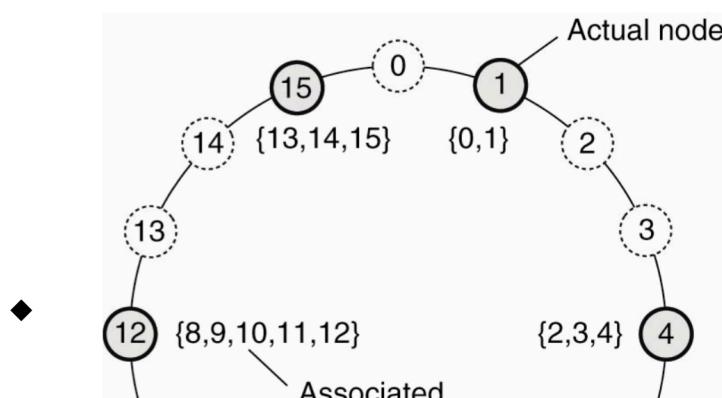
e

se





- Example:
  - Cloud computing:
    - ◆ Cloud computing: the delivery of computation or storage as a service
    - ◆ The client hosts the user interface to launch the computation
    - ◆ The servers handle most of the computation upon request and return the results
    - ◆ One server asks data to another server
    - ◆ Another does the computation
- Peer-to-peer organization:
  - P2P model:
    - Every machine acts similarly
      - Every machine is both a client and a server
      - Not centralized control: the responsibility is distributed evenly
      - Even the program executing on each machine is similar
    - Examples:
      - Chord, a Distributed Hash Table (DHT)
        - ◆ As a node:
          - ◊ Have a successor peer
          - ◊ Have a predecessor peer
          - ◊ Have some shortcuts to other nodes to speedup delivery
          - ◊ Be responsible of a subset of the system data items (base range)

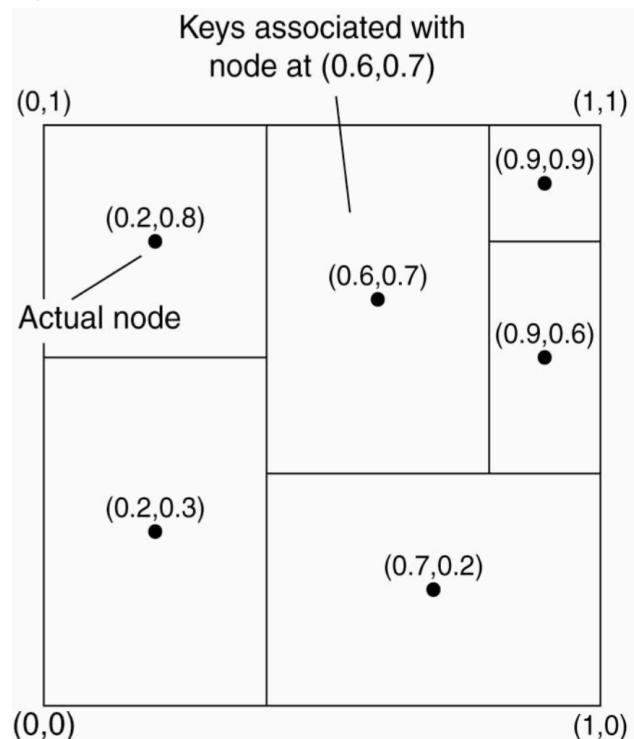


service to end-users  
and prints the results  
d sends back the results to the client:

y of requests  
sed on my unique identifier)



- Associated data keys
- 
- Content Addressable Network (CAN), another DHT
    - ◆ As a node:
      - ◊ Be responsible of a region of the system (based on my own IP)
      - ◊ Have few neighbors, the nodes with adjacent regions, who are responsible for the adjacent regions



- BitTorrent, a file sharing application
  - ◆ Used for LINUX distribution, software patches, distributing files
  - ◆ Goal: quickly replicate large files to large number of clients
  - ◆ Web server hosts a .torrent file
  - ◆ A tracker (server or a DHT) tracks downloaders/ owners of a file
  - ◆ Files are divided into chunks
  - ◆ Downloaders download chunks from themselves (and owner)
  - ◆ Tit-for-tat: the more one shares (server), the faster the it can download

- Distributed operating system:
  - Distributed operating system:
    - A single system image, the system maintains a single copy of the resources



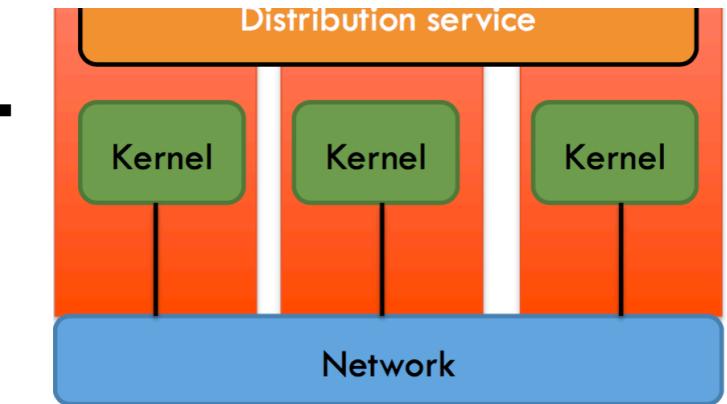
unique identifier)  
with which I can communicate

ile

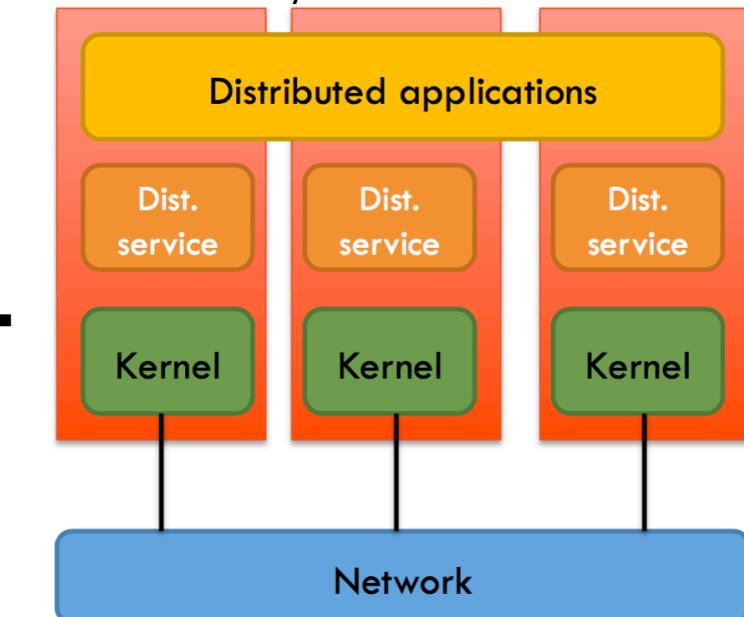
download (client)

s

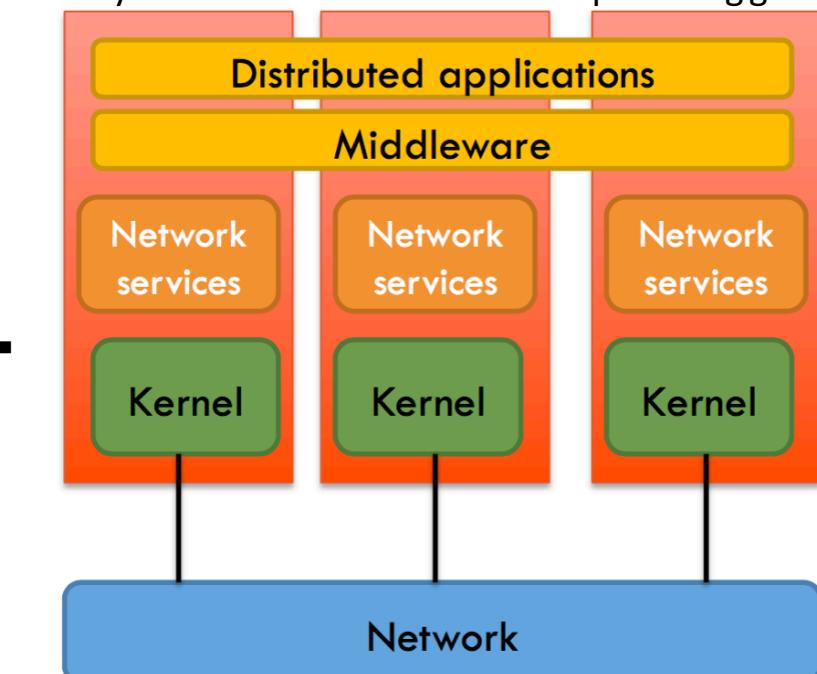




- o Network operating system:
  - Machines provide resources to other machines (e.g., UNIX rlogin)
  - The OS can vary from one machine to another, essentially file sharing



- o Middleware:
  - A layer over the network services providing general services to application

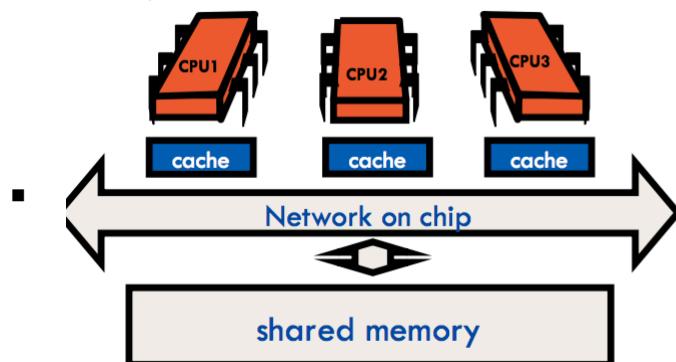


ns in a vary transparent manner (system can differ)



## Concurrency

- Uniprogramming and multiprogramming
  - o Uniprogramming: one program execution at a time
    - MS/DOS
    - Early macintosh
    - Batch processing
  - o Many program execution on personal computers:
    - Browsing the web
    - Sending emails
    - Typing a letter
    - Burning a DVD
  - o Multiprogramming: more than one program execution at a time
    - Multics
    - UNIX/Linux
    - Windows NT
    - Mac OS X
  - o Multi/many-cross



- Concurrency:
  - o Assume a single central processing unit (CPU)
  - o The way to give the illusion to the users of multiple CPUs:
    - Answer: Scheduling program execution, one after the other during some time

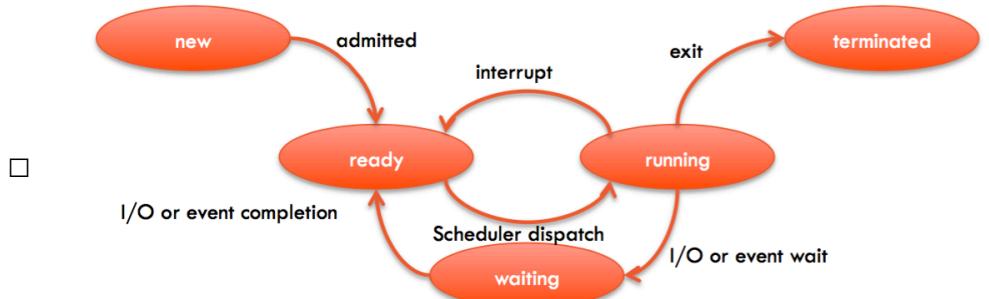


## Processes and threads:

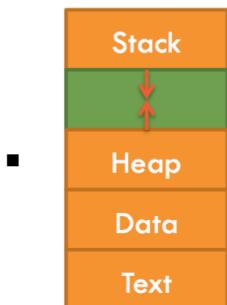
- Process:
  - o An abstraction representing a program execution

small fractions of

- When to switch from one process to another?
  - Timer interrupt goes off, explicit yield, I/O, etc.
- Each virtual CPU needs to store the process context:
  - Program counter, stack pointer, other registers
- A process switches from a state to another:



- Address space:
  - A unit of management of a process's virtual memory
  - Process address space:
    - Stack: temporary data extensible towards lower virtual addresses
    - Heap: memory allocated dynamically extensible to higher virtual addresses
    - Data section: global variable
    - Text region: program code



- Process management:
  - Process creation:
    - Process allocation: the act of choosing the host of the process
      - The system V ("5") provide command to execute a process at another workstation
      - The amoeba system choose a host from a pool of processors to run the process
    - Execution environment: an address space w/ initialized content and operating system support
      - Static: program text, heap and stack regions created from a list of segments
      - Dynamic: UNIX fork shares program text and copies stack and heap
  - Interprocess communication (IPC):
    - Pipes/message queue/shared memory segments
    - Requires costly context switches
- Context-switch:
  - An interrupt requires kernel intervention
  - S1. process a switches from user to kernel mode
    - Changing the memory map in the MMU / memory management

esses

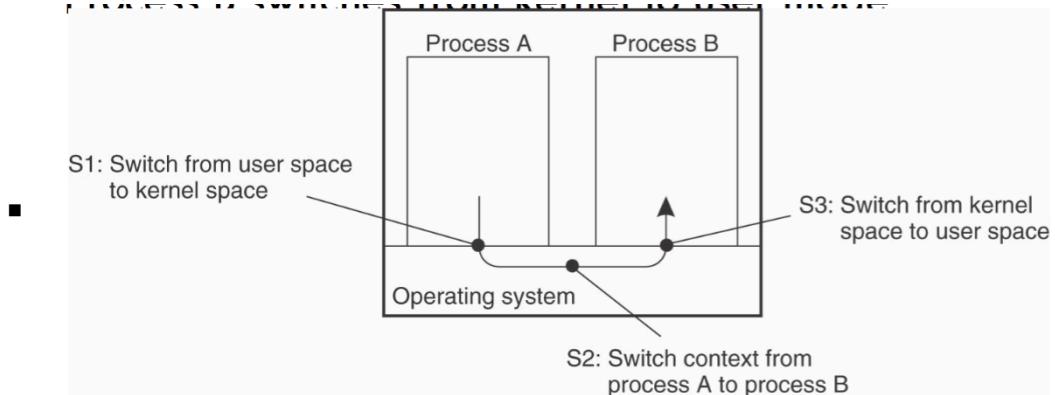
idle

execute it  
open files:

heap regions

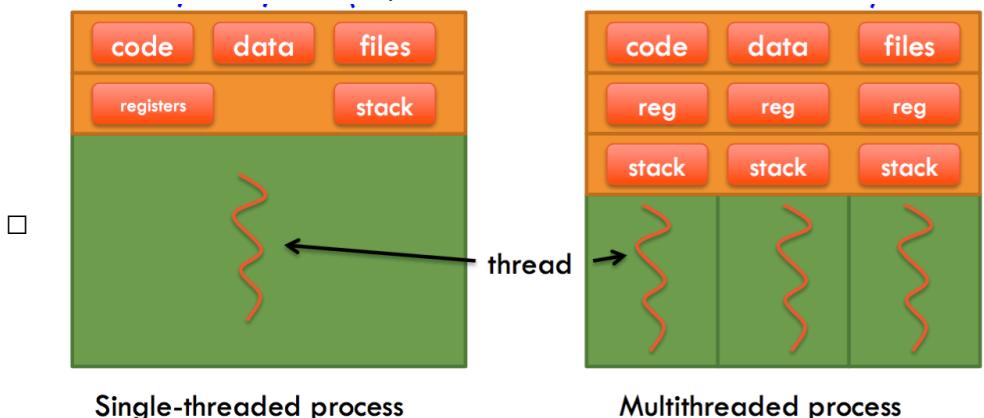
unit)

- □ Changing the memory map in the MMU (memory management unit)
- □ Flushing the TLB (translation lookaside buffer)
- □ S2. context-switch (swapping processes from disk to main memory)
- □ S3. process B switches from kernel to user mode



- Thread:

- Smallest unit of CPU utilization
  - Portable OS interface (POSIX) threads
    - Current activity: program counter
    - Stack: temporary data
    - No data, code, files (but it shares them with other threads)



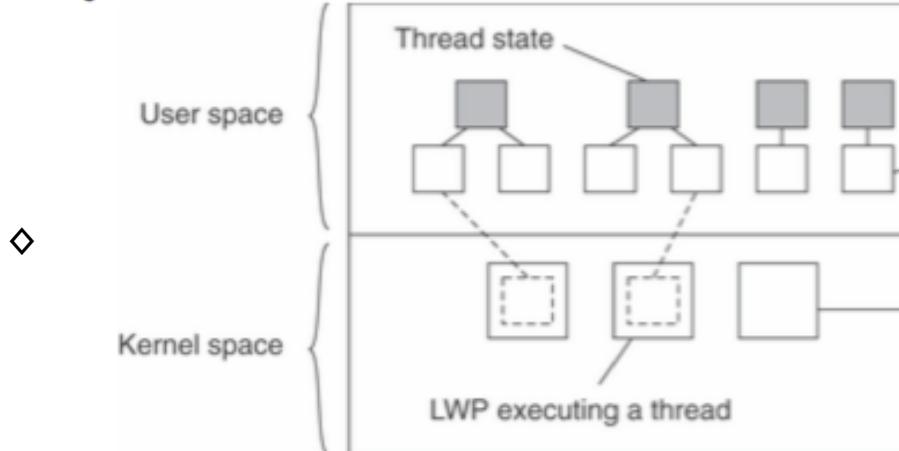
- Communication between threads always through memory
  - Does not need IPC
  - No context switches required (when purely at user-level)
- User-level threads library and kernel-scheduled thread
  - User-level thread library:
    - ◆ Cheap to create and destroy
    - ◆ Blocking system call freezes the entire process of the threads)
    - ◆ Cheap context-switch:
      - ◊ Few instructions to store and reload CPU register values
      - ◊ No need to change memory map in MMU
      - ◊ No need to flush the TLB
  - Kernel-scheduled threads:

urn:

ad (and other

ues

- ◆ Costly to create and destroy
  - ◆ Do not block the current process (and other threads) upon system call (I/O)
  - ◆ Costly context-switch (similar to process context switch)
    - ◊ Needs to change memory map in MMU
    - ◊ Needs to flush the TLB
- User-level threads library and kernel-scheduled thread
  - Lightweight processes (LWP):
    - ◆ Runs in the context of a process (potentially multiple LWP)
    - ◆ Current threads are in a table
    - ◆ Once an LWP finds a runnable thread:
      - ◊ It locks the table in user-mode to update it
      - ◊ It switches context to that thread in user-mode
    - ◆ Does not block upon system calls:
      - ◊ Passing from user to kernel mode in the LWP context
      - ◊ If the LWP is blocked, the kernel context-switches to another LWP
      - ◊ Implying a context-switch also at the user level



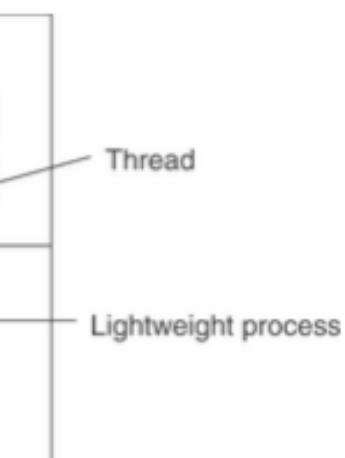
- Process and thread:
  - Process:
    - ◆ Isolated: prevents one process from interfering with another
    - ◆ Inefficient: starting/ terminating a process and context switch
  - Threads:
    - ◆ Non-isolated: avoiding incorrect interferences makes programs more reliable
    - ◆ Efficient: a thread is a lightweight version of a process
- Client multi-threading
  - Ex: Google Chrome web browser
    - ◆ Multithreading: each tab runs its own thread
    - ◆ A JavaScript error does not crash the main Chrome process
      - ◊ Only one tab freezes, the user can close it independently
    - ◆ Takes benefit of multicore architectures, each thread runs on a core

in blocking

s per process)

xt

o another LWP



ther

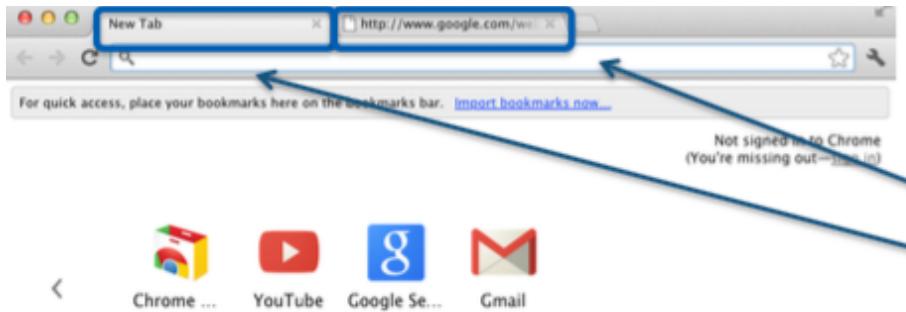
switches are costly

gramming harder

s:

ently of others

s on a separate



- Java threads:



one distinct thread  
running each tab