

COMP2121: Principles of Distributed Systems and Networks

Lecture title

Unit coordinator Dr. Vincent Gramoli
School of Information Technologies



THE UNIVERSITY OF
SYDNEY



Motivations

- Entities are distributed so they may act concurrently to update some data
- Data may also be distributed (i.e., replicated) and not all replicas can be accessed simultaneously
- What does it mean to be consistent?

Outline

- Definitions
- Sequential consistency
- Causal consistency
- Replication
- Multiple-access operations

Definitions



THE UNIVERSITY OF
SYDNEY

Definitions

Problem of inconsistency

- Data is distributed (replicated) to
 - Improve the **reliability** of a system
 - **Scale** in numbers: preserve performance while multiplying entities
 - Scale in geographical area: preserve performance while increasing space
- Drawback: **cost** of keeping the replicas consistent
 - All replicas must have to maintain information
 - Intuitively, a client update must be propagated to **all** replicas
 - but a client can obtain an “up-to-date” information from **any** replica
- Not so simple!
 - Propagation takes time and one replica may be up-to-date while another is not
 - *Inconsistencies*: two clients may observe a different data depending on the replica they contact

Definitions


Conflict

- Two (groups of) operations are *concurrent* in the traditional sense here:
 - Each of the two (groups of) operations starts before the other ends
- An *update operation* is an operation that modifies the data
 - E.g., a write operation or a transaction (i.e., a group of operations) with a write
- A *conflict* is a relation between two (groups of) operations:
 - that are concurrent,
 - access the same data, and
 - (at least) one of which is an update

Definitions

Data store

- A *data store* is a distributed collection of storages



- From the standpoint of a given process, only one storage copy represents its *local copy*

Definitions

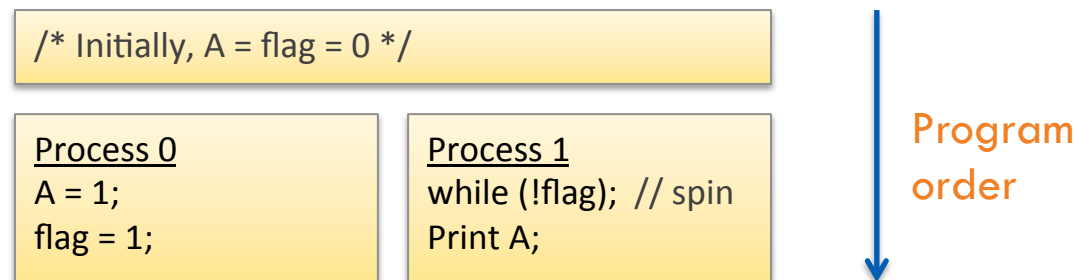
Consistency

- A *consistency model* is a **contract** between the processes and the data store such that **if processes obey certain rules, the store promises to work correctly**
- Data-centric vs. client-centric consistency
 - Data-centric consistency: consistency in the general sense (e.g., causal consistency, sequential consistency)
 - Client-centric consistency: consistency for a single client (independently from other clients), assuming that a single client does not trigger multiple concurrent updates (e.g., eventual consistency)

Definitions

Coherence vs. consistency

- **Coherence** concerns only **one memory location**
- **Consistency** concerns apparent ordering for **all locations**



- Example: output “0” is coherent but non (sequentially) consistent
- A memory system is coherent if for **each** location, we can serialize **all operations to that location** such that, operations performed by any processor appear in **program order** (i.e., the value returned by a read is the value written by the last write to that location)

Sequential Consistency



THE UNIVERSITY OF
SYDNEY

Sequential consistency

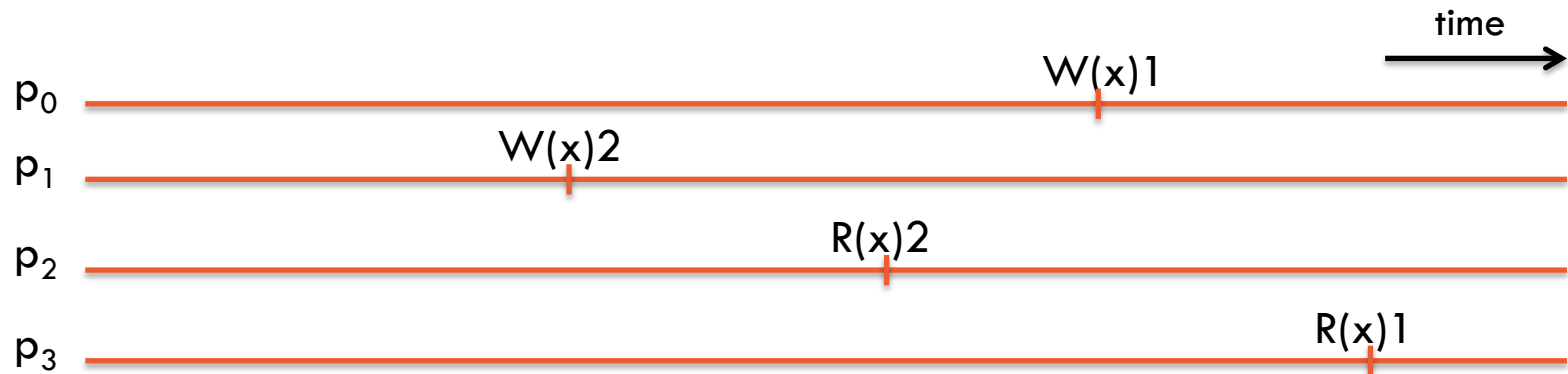
Distributed execution

- *Program order*: The order in which events appear to be executed **locally** (in the program)
 - *Execution*: A sequence of **reads** (R) and **writes** (W) on some **data items** $x \dots$ executed by **distributed processes** $p_0 \dots$
 - *Serial execution*: An execution where each process executes (all its operations) **one after another**
 - *Equivalent executions*: Two executions are equivalent if they contain exactly the same reads and writes:
 - executed from the same process,
 - all writes have the same input values and
 - all reads have the same output values
 - Notation:
 - $W_i(x)1$ is a write operation by process p_i on data item x with value 1 (index i omitted when clear from context)
 - $R_i(x)2$ is a read operation by process p_i on data item x that returns value 2 (index i omitted when clear from context)
-
- The diagram illustrates a distributed execution timeline for four processes: p_0 , p_1 , p_2 , and p_3 . Each process is represented by a horizontal orange line. Vertical red tick marks indicate the execution of specific operations. The operations are: $W(x)2$ by p_1 , $W(x)1$ by p_0 , $R(x)2$ by p_2 , and $R(x)1$ by p_3 . A black arrow labeled "time" points to the right, indicating the progression of time. The operations are ordered sequentially by time: $W(x)2$ occurs first, followed by $W(x)1$, then $R(x)2$, and finally $R(x)1$.
- Example 1: the execution $W_1(x)2; R_2(x)2; W_0(x)1; R_3(x)1$ is **serial** in that each process executes one after another in the order $p_1; p_2; p_0; p_3$

Sequential consistency

Sequentially consistent executions

- *Sequential consistency*. The result of each execution is
 - the same **as if** the (read and write) operations by all processes on the same data store were executed in **some sequential order** and the operations of each individual process appear in this sequence in **its program order**

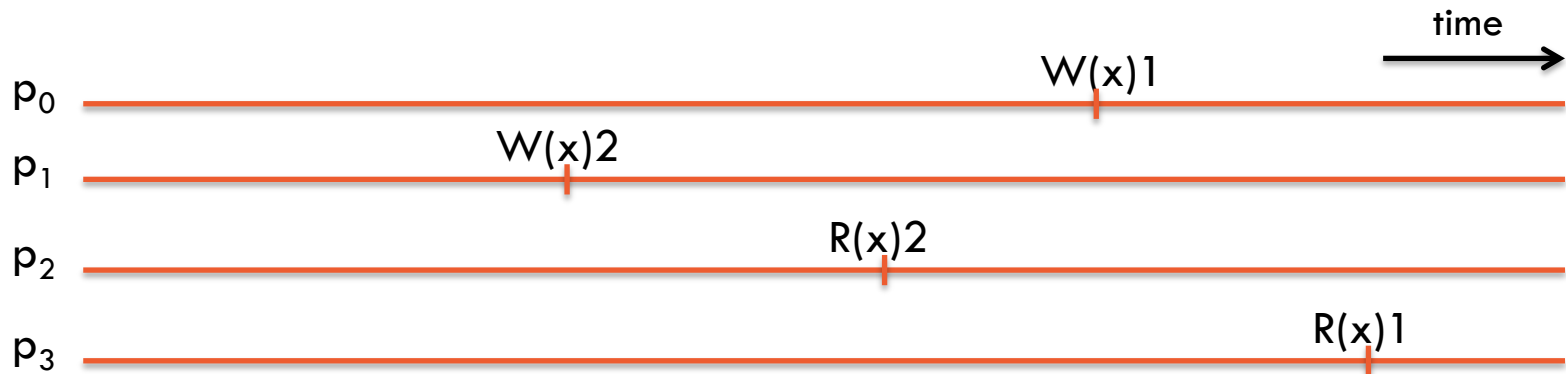


- Example 1 (bis): is this execution sequentially consistent as well?

Sequential consistency

Sequentially consistent executions

- *Sequential consistency.* The result of each execution is
 - the same as if the (read and write) operations by all processes on the same data store were executed in some sequential order and the operations of each individual process appear in this sequence in its program order




- Example 1(bis): the result of this serial execution is **sequentially consistent** as well

If an execution is serial, then its result is always sequentially consistent

Sequential consistency

A non-sequentially consistent execution

- *Sequential consistency.* The result of each execution is
 - the same as if the (read and write) operations by all processes on the same data store were executed in some sequential order and the operations of each individual process appear in this sequence in its program order




- Example 2: the result of this execution is **sequentially consistent**
 - The result of the execution is $W_0(x)1; W_1(x)2; R_2(x)2; R_3(x)2; R_2(x)1; R_3(x)1$
 - An equivalent sequential order with the same result is $W_1(x)2; R_2(x)2; R_3(x)2; W_0(x)1; R_2(x)1; R_3(x)1$
 - The program order of each process is satisfied by this sequence:
 - $R_3(x)2$ before $R_3(x)1$ and
 - $R_2(x)2$ before $R_2(x)1$

Sequential consistency

A non-sequentially consistent execution

- *Sequential consistency.* The result of each execution is
 - the same as if the (read and write) operations by all processes on the same data store were executed in some sequential order and the operations of each individual process appear in this sequence in its program order




- Example 3: the result of this execution is **not sequentially consistent**
 - The result of the execution is $W_0(x)1; W_1(x)2; R_2(x)2; R_3(x)1; R_2(x)1; R_3(x)2$
 - There is no sequential execution in which:
 - By program order of p_3 , $R_3(x)1$ is before $R_3(x)2$ (implying that $W_0(x)1$ would be before $W_1(x)2$ and
 - By program order of p_2 , $R_2(x)2$ before $R_2(x)1$ (implying that $W_1(x)2$ would be before $W_0(x)1$)

⇒ **contradiction**


Sequential consistency

Application to shared memory multiprocessors

- How can a coherent execution may not be sequentially consistent



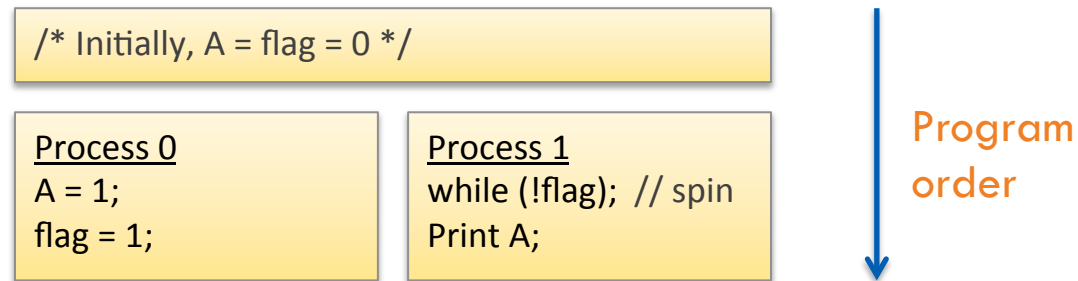
- Example 4: What are the possible outputs of this concurrent program?




Sequential consistency

Application to shared memory multiprocessors

- How can a coherent execution may not be sequentially consistent




- Example 4: What are the possible outputs of this concurrent program?




Sequential consistency

Application to shared memory multiprocessors

- How can a coherent execution may not be sequentially consistent



- Example 4: output “0” is not (sequentially) consistent as we cannot find a sequential order that respects program order and prints 0.



Causal Consistency

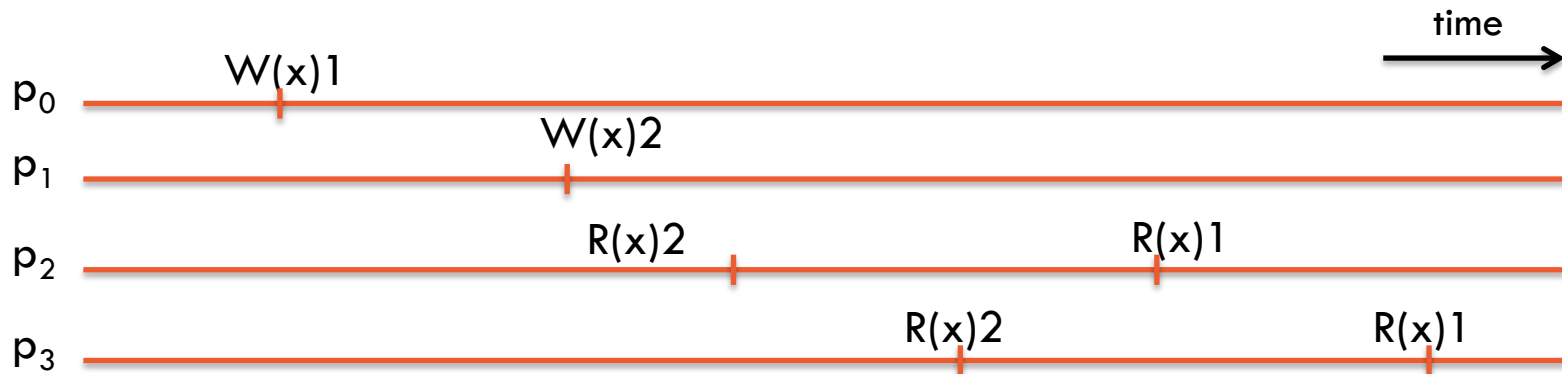


THE UNIVERSITY OF
SYDNEY

Causal consistency

Causally consistent executions

- *Causal consistency*. Writes that are *causally related* (i.e., one happens before the other) must be seen by all processes in the same order. Concurrent writes may be seen in a different order on different machines




- Example 2(bis): is this execution causally consistent?

Causal consistency

Causally consistent executions (con't)

- *Causal consistency.* Writes that are *causally related* (i.e., one happens before the other) must be seen by all processes in the same order. Concurrent writes may be seen in a different order on different machines




- Example 2(bis): is this execution causally consistent?
 - Write to the same data items are not necessarily causally related
 - A read returning the value written are causally related

Causal consistency

Causally consistent executions (con't)

- *Causal consistency*. Writes that are *causally related* (i.e., one happens before the other) must be seen by all processes in the same order. Concurrent writes may be seen in a different order on different machines




- Example 2(bis): the result of this serial execution is **causally consistent**

If the result of an execution is sequentially consistent, then it is also causally consistent

Causal consistency

Causally consistent executions (con't)

- *Causal consistency*. Writes that are *causally related* (i.e., one happens before the other) must be seen by all processes in the same order. Concurrent writes may be seen in a different order on different machines

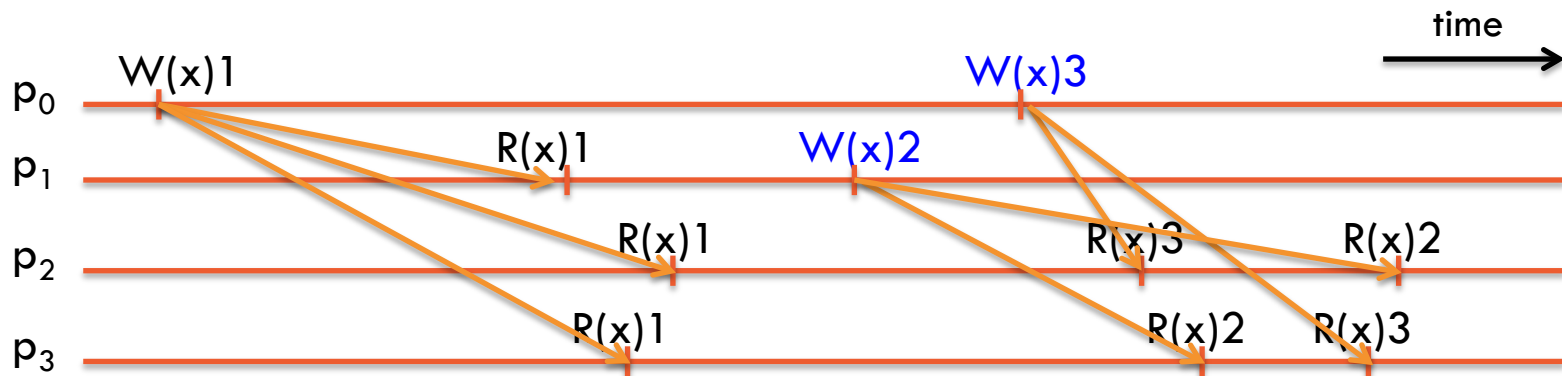


- Example 5: is the result of this execution is causally consistent?

Causal consistency

Causally consistent executions (con't)

- *Causal consistency*. Writes that are *causally related* (i.e., one happens before the other) must be seen by all processes in the same order. Concurrent writes may be seen in a different order on different machines




- Example 5: the result of this execution is **causally consistent**
 - $W_0(x)1$ and $W_0(x)3$ that are causally related are seen in the same order by p_0 , p_1 , p_2 and p_3
 - $W_1(x)2$ and $W_0(x)3$ are **not causally related** as none happen before the other

Vector clocks can be used to implement a causally consistent data store

Causal consistency

Causally consistent executions (con't)

- *Causal consistency.* Writes that are *causally related* (i.e., one happens before the other) must be seen by all processes in the same order. Concurrent writes may be seen in a different order on different machines




- Example 5: but is it sequentially consistent?

Causal consistency

Causally consistent executions (con't)

- *Causal consistency*. Writes that are *causally related* (i.e., one happens before the other) must be seen by all processes in the same order. Concurrent writes may be seen in a different order on different machines

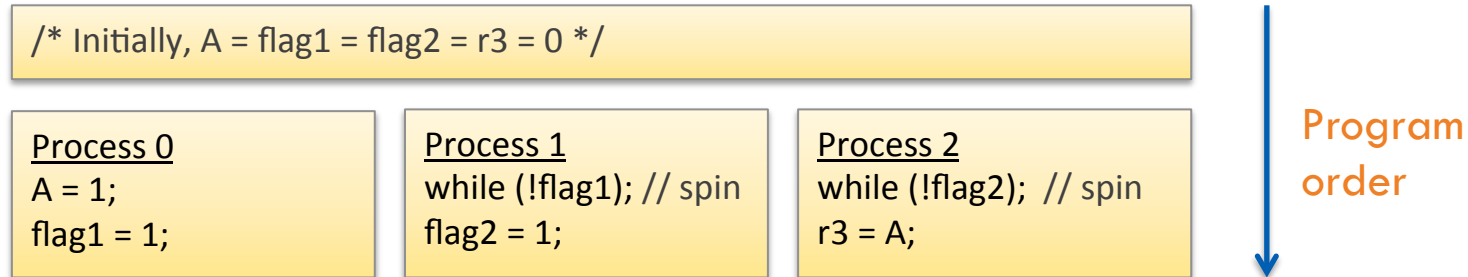


- Example 5: the result of this execution is **not sequentially consistent**


Causal consistency

Application to shared memory multiprocessors

- All commercialized memory models guarantee causal consistency



- Causal consistency implies that at the end $r3 = 1$, because all the writes are causally related (we can order them with the *happen-before* relation)



Replication




THE UNIVERSITY OF
SYDNEY

Replication

Updating a replicated system

- Writes operations can be carried out to several sites (like in the database example)



- Updates must be executed in **the same order** at each replicated site
- The use of **totally-ordered broadcast** (cf., previous lecture) can solve this
- However, our previous Lamport's clock solution **does not scale well**

Replication

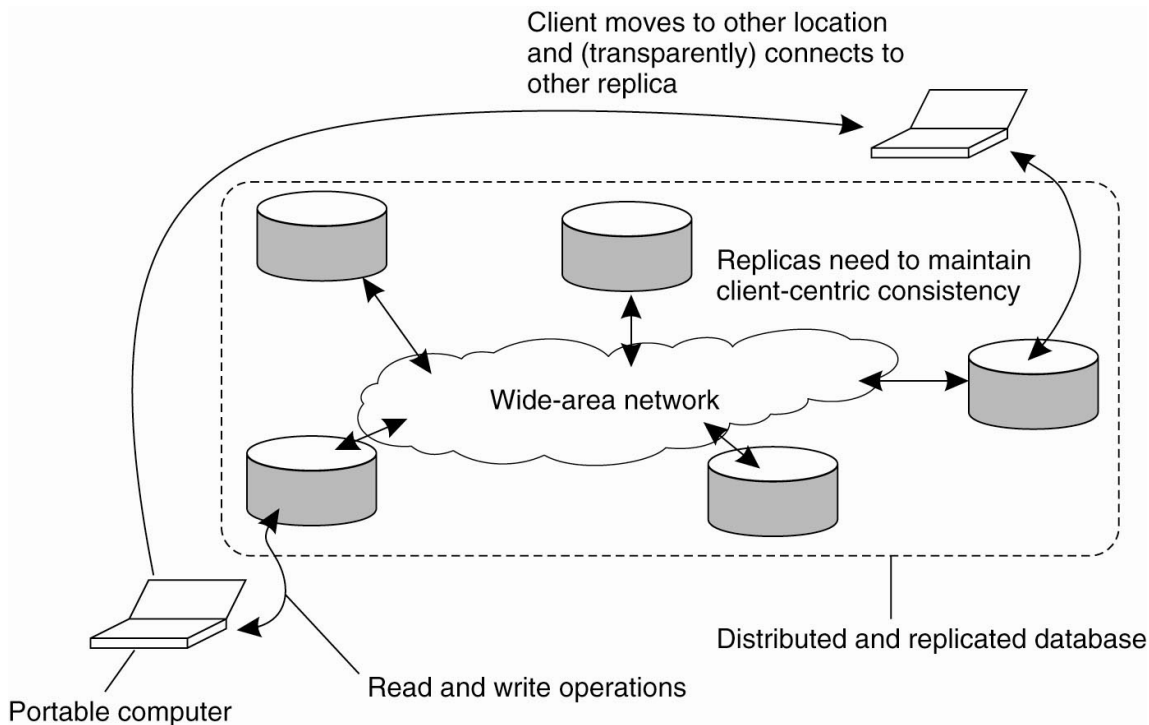
Eventual consistency: eventually all replicas are consistent

- Assumption: **no concurrent updates**
 - The only type of conflicts occur between a reading operation and an updating operation (read-write conflict)
 - Never between two updating operations (no write-write conflicts)
- Example:
 - A web server is the only one to update the webpage content
 - Clients typically access the webpage in read mode only (without modifying it)

Replication

Eventual consistency (con't)

- A mobile user accesses a database by connecting to one of its replicas in a transparent way (the user's application is unaware of which replica)
- The user executes several update requests and disconnects
- The user then reconnects from a different location or through a different device.
- If the changes have not been propagated by the system to all replicas, the user may observe a client-centric **inconsistency**



Replication

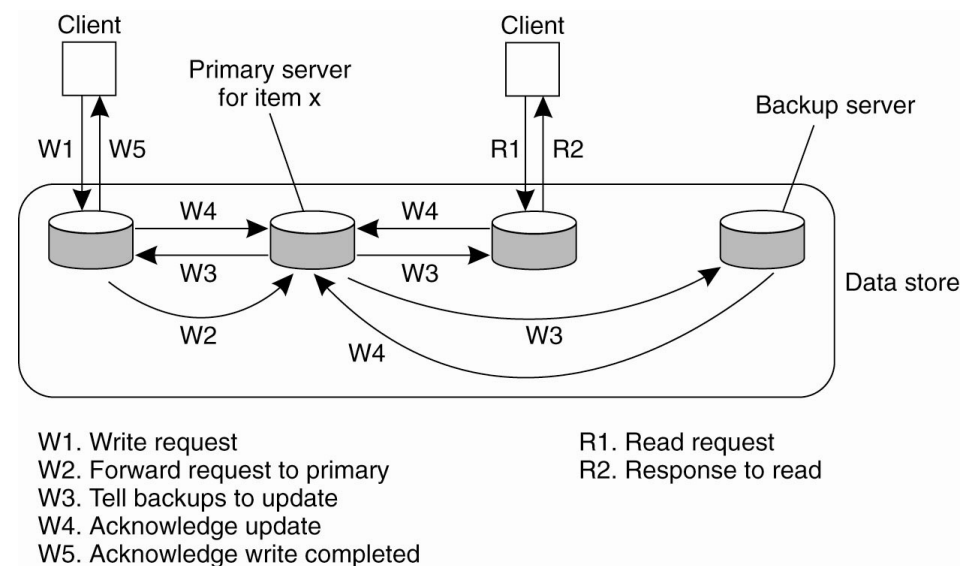
Coordinator

- Active replication
 - Forward **all messages to a central coordinator (leader)**
 - The coordinator is a *sequencer* that chooses a **unique sequence number for each message**
 - It sends this sequence number along with the message **to all replicas**
 - Operations are carried out in the **order of their sequence number**

Replication

Primary-backup

- Primary-backup protocol [Budhiuraja et al., 1993]
 - All write operations on x are forwarded to a **primary server** for x
 - The **primary does the update** and forwards the request to **backup servers**
 - Each backup executes the updates and **acknowledges the primary**
 - The primary sends a **response to the client**
 - Read operations can be carried out **locally**
- Pro: **guarantees sequential consistency**
- Cons: **delayed answer to client due to blocking execution**

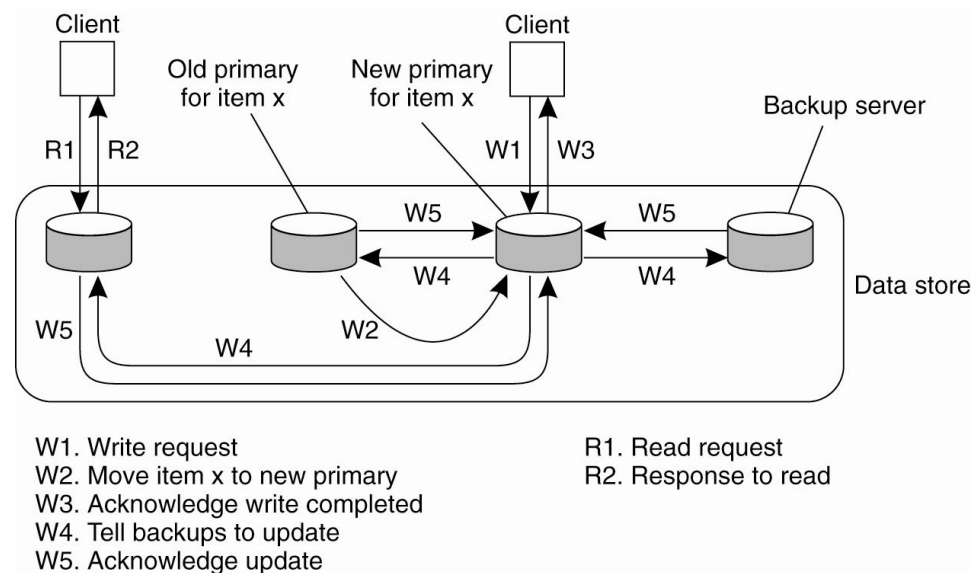


Replication

Primary-backup

– Local-write primary-backup

- When a process wants to update a data item x **it locates it**
- It **moves the item** to its location, becoming the new primary and updates the item locally
- It **acknowledges** the client
- It **tells the backup** to execute the update locally
- The backups acknowledge the new primary
- Pro: **does not block**



Replication

Quorum system [Gifford 1979, Thomas 1979]

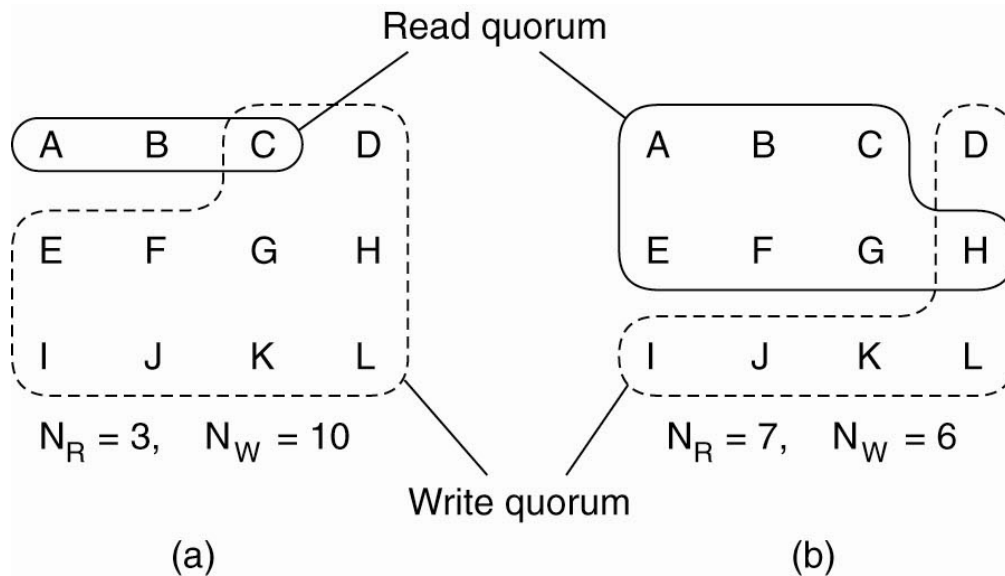
- *Quorum system*: a set S of quorums that are mutually intersecting sets
- Property: for any pair of quorums $Q1, Q2$ in the quorum system S we have:
 $Q1 \cap Q2 \neq \emptyset$.
- Key idea:
 - Consider a **single quorum system**
 - Write:
 - forward the update request along with a **version** that is a pair of timestamp and node identifier, to (all replicas of) a **quorum**
 - Once **acknowledged by a quorum** of replicas, **respond** to client
 - Read:
 - When a **read** on x is requested, forward the request to all replicas belonging to a quorum and picks the response with the maximum timestamp
 - Return the **chosen copy** to the client
- **Nodes have to know about the quorum system**

Replication


Quorum system [Gifford 1979, Thomas 1979]

- Gifford's approach: Consider N replicas.
 - Upon read request, forward the request to at least N_R servers (a *read quorum*)
 - Upon writing, forward the request to at least N_W servers (a *write quorum*) such that:

$N_R + N_W > N$ (each *read quorum intersects* with each *write quorum*) and
 $N_W > N/2$ (each *write quorum intersects* with each *write quorum*)



(b) Cannot work as $N_W = N/2$



Multiple-Access Operations



THE UNIVERSITY OF
SYDNEY

Multi-access operations


Definition

- Any **single** operation can now access **multiple** objects (e.g., *transaction*)
- We have to consider **new consistency criteria** (previous ones do not apply)

Primitive	Description
BEGIN_TRANSACTION	Mark the start of a transaction
END_TRANSACTION	Terminate the transaction and try to commit
ABORT_TRANSACTION	Kill the transaction and restore the old values
READ	Read data from a file, a table, or otherwise
WRITE	Write data to a file, a table, or otherwise

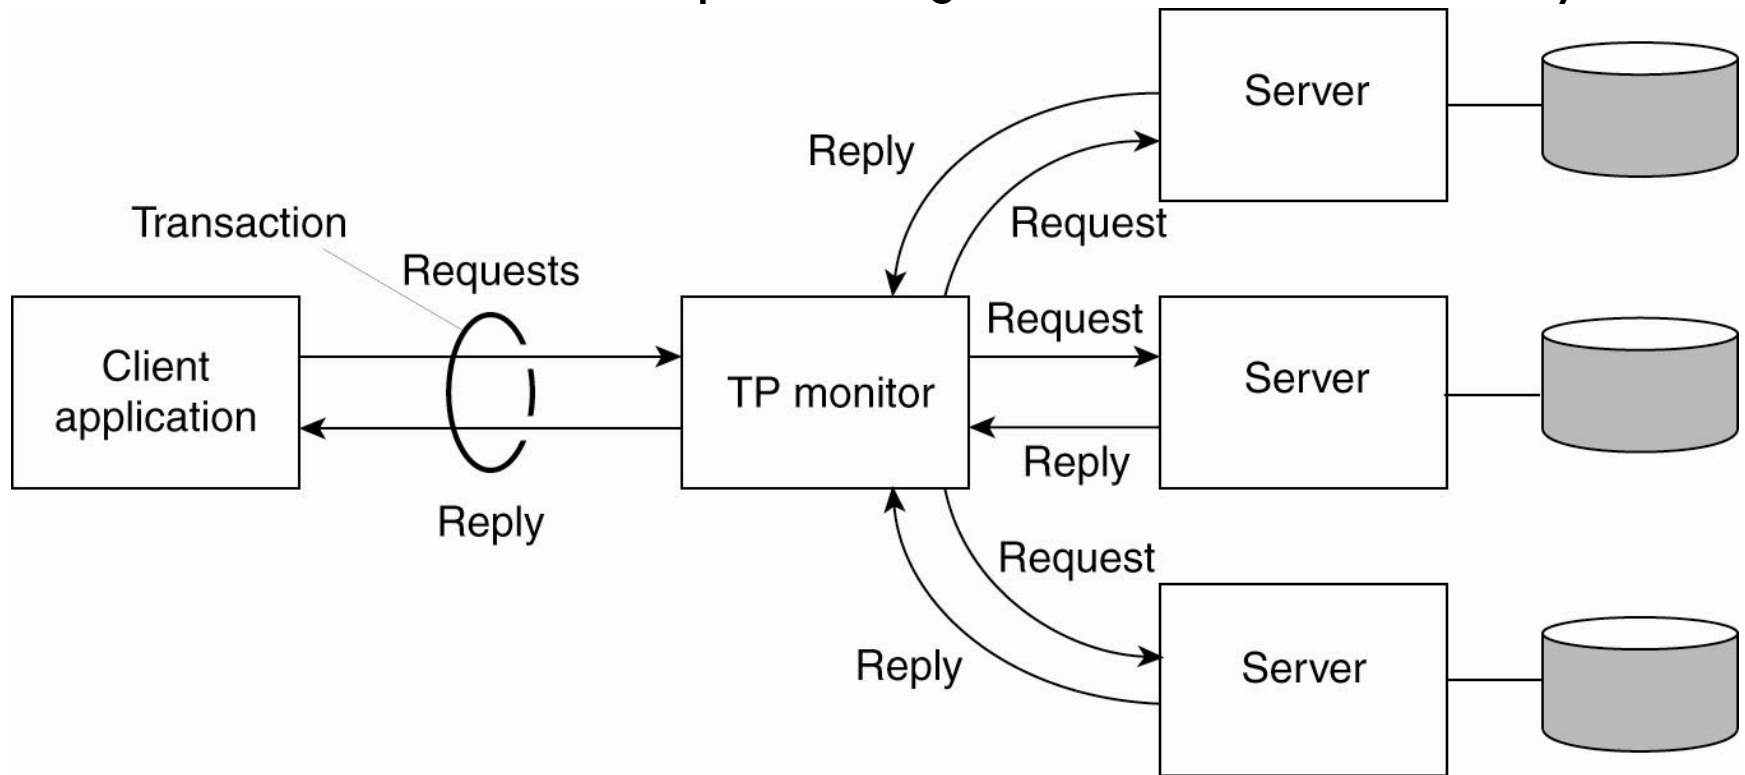
Database transactions

- Nested transactions



Database transactions

- The role of a transaction processing monitor in distributed systems



Database transactions

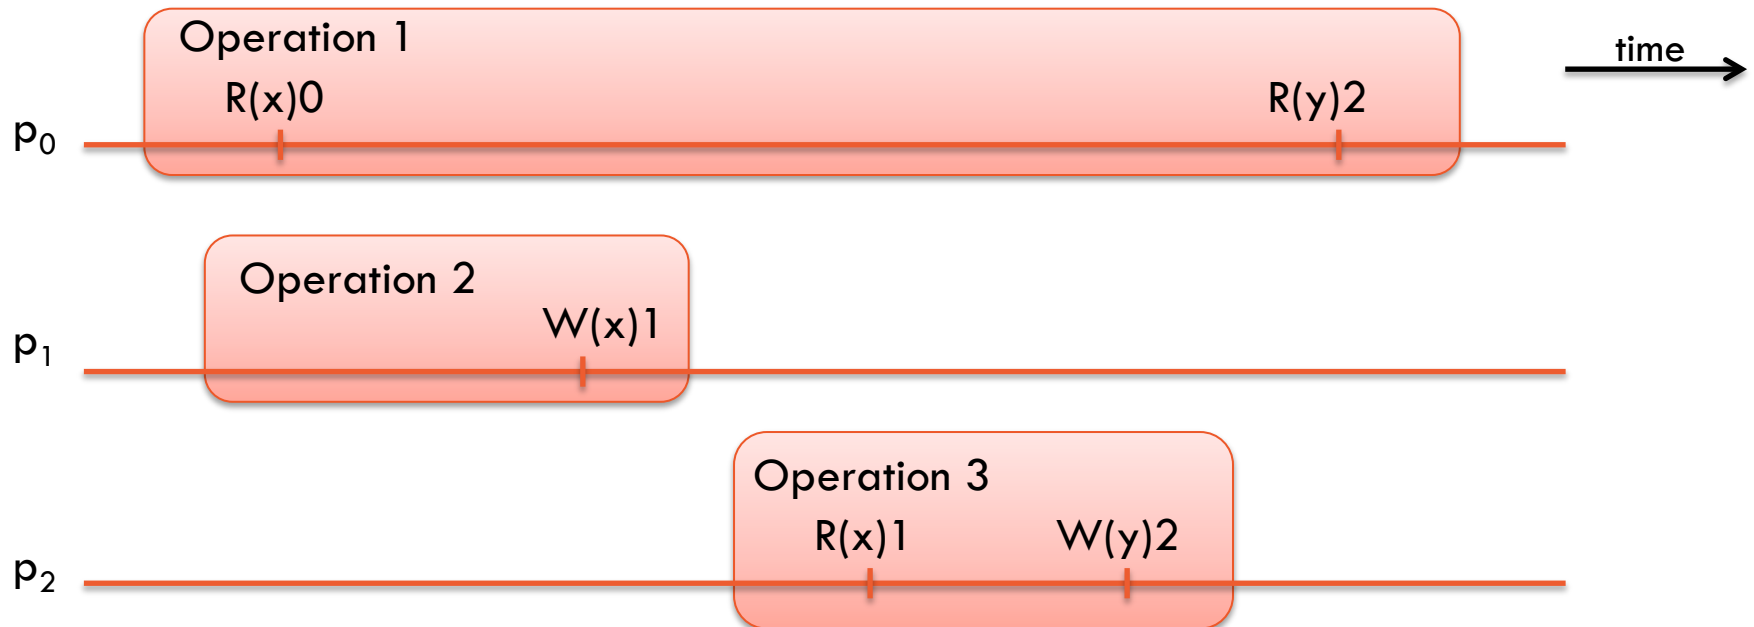
Properties

- Characteristic properties of transactions:
 - Attomic: To the outside world, the transaction happens indivisibly
 - Consistent: The transaction does not violate system invariants
 - Isolated: Concurrent transactions do not interfere with each other
 - Durable: Once a transaction commits, the changes are permanent
- These are named the *ACID* properties

Serializability

Serializability [Papadimitriou, JACM 1979]

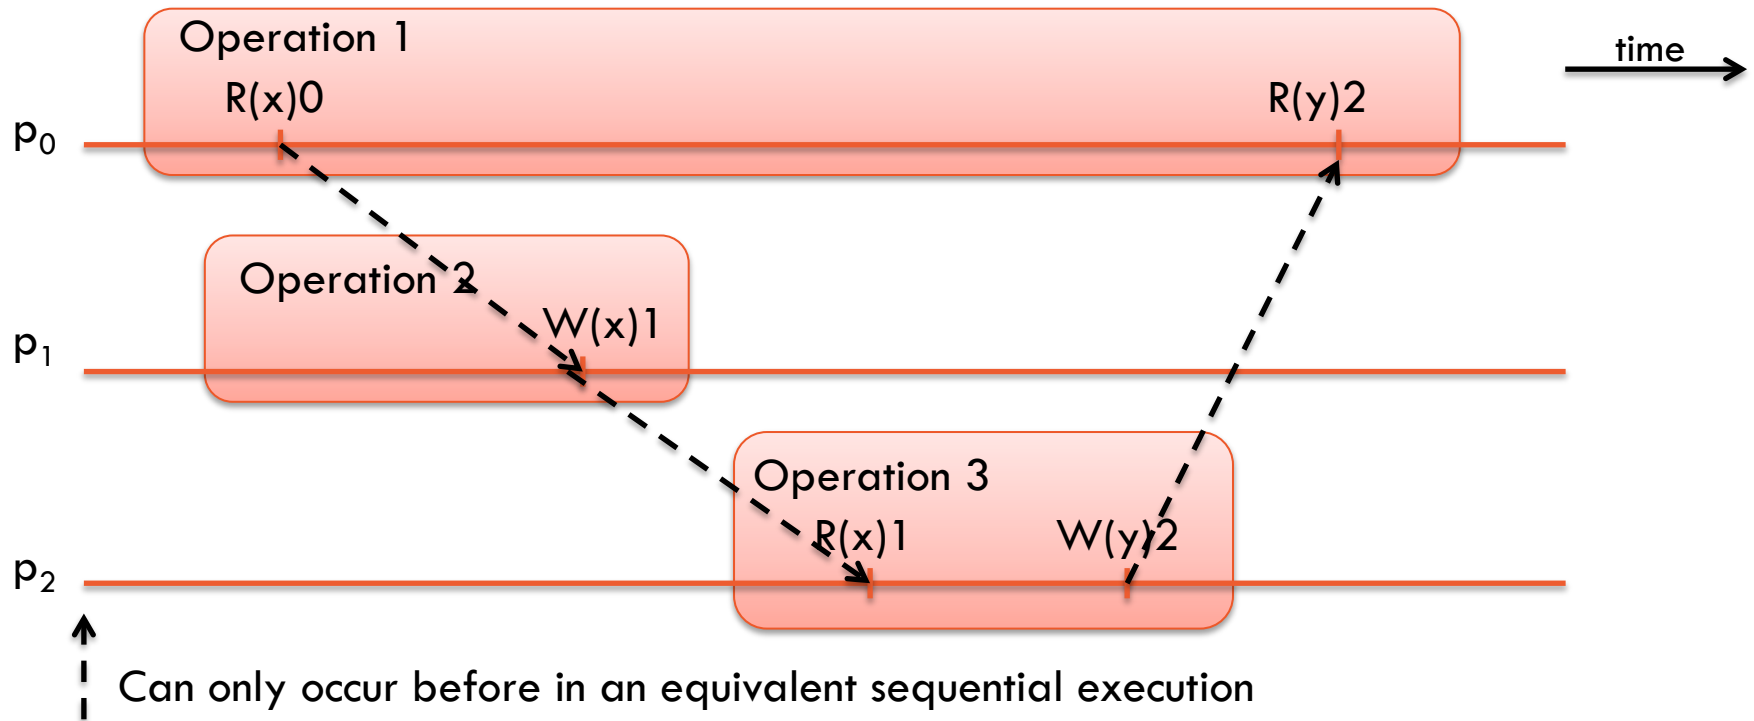
- *Serializability*: the result of an execution (of multi-access operations) is serializable if there exists an equivalent sequential execution.



Serializability

Serializability

- *Serializability*: the result of an execution (of multi-access operations) is serializable if there exists an equivalent sequential execution.

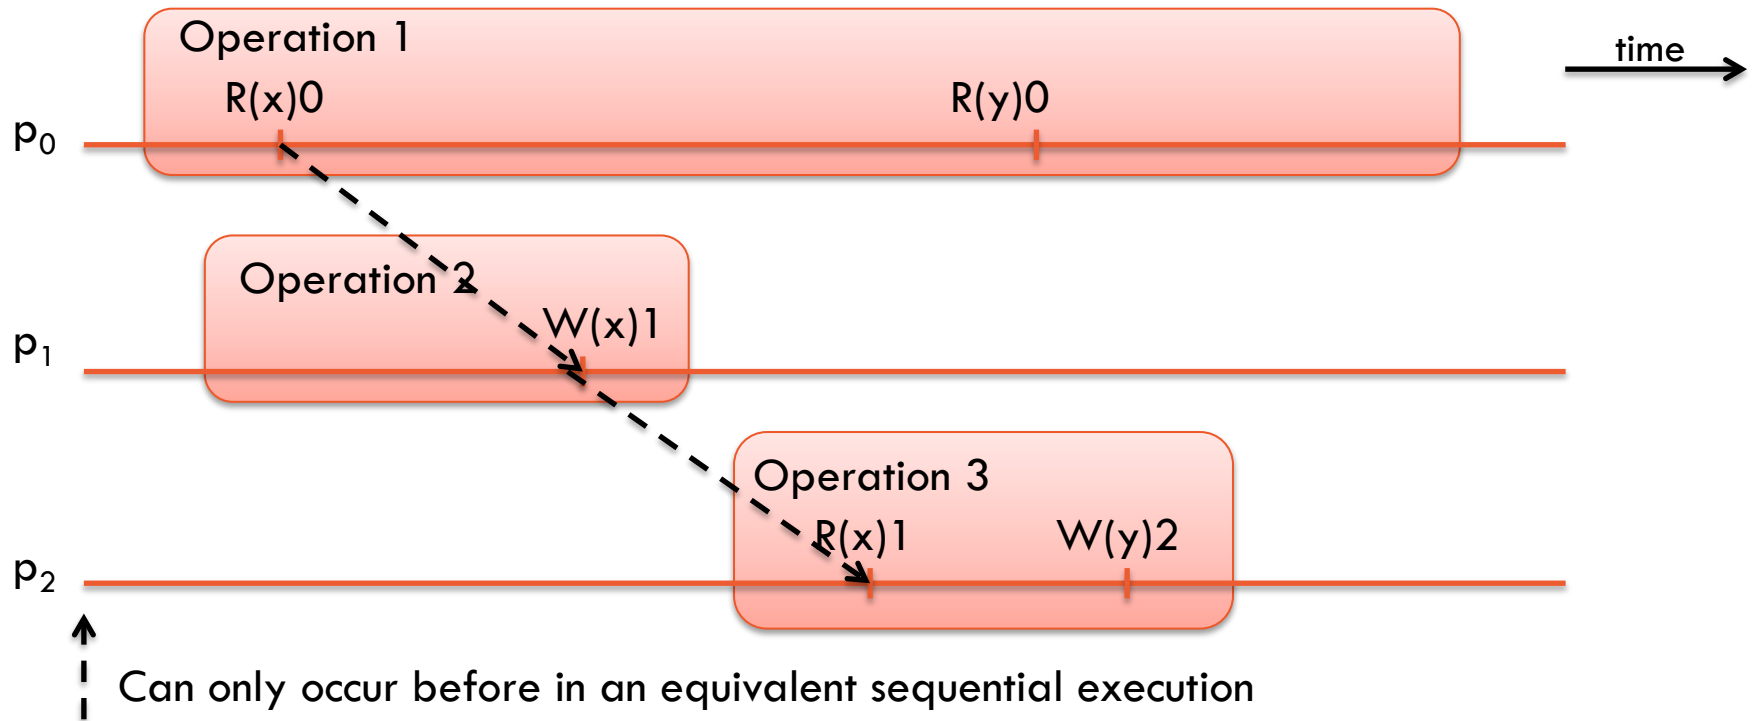


There is a cycle $op1 \rightarrow op2 \rightarrow op3 \rightarrow op1$ in the precedence graph \Rightarrow non-serializable

Serializability

Serializability

- *Serializability*: the result of an execution (of multi-access operations) is serializable if there exists an equivalent sequential execution.

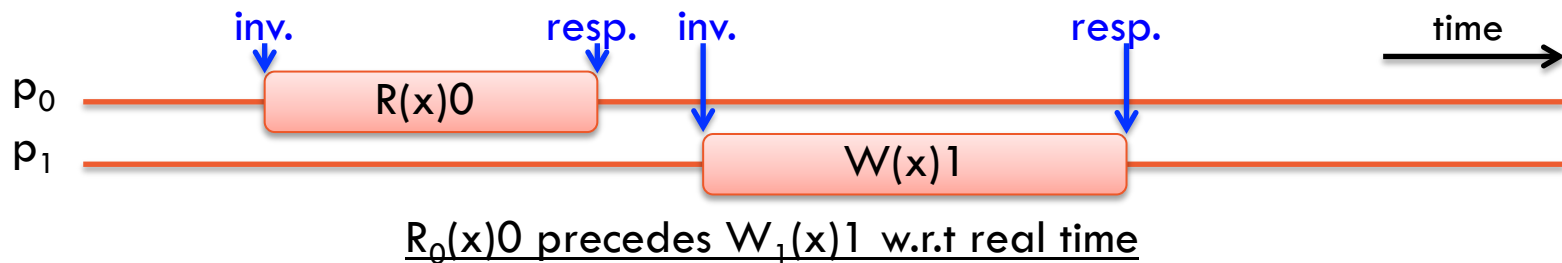


No cycle $op1 \rightarrow op2 \rightarrow op3$ in the precedence graph \Rightarrow serializable

Linearizability

Linearizability (of reads and writes) [Herlihy and Wing, TOPLAS 1990]

- *Real-time precedence*: if an operation o_1 returns before another operation o_2 is invoked (at a potentially different process) then o_1 *precedes* o_2 *with respect (w.r.t) to real-time*.




- *Linearizability*. The result of each execution is
 - equivalent to a sequential execution that respects the real-time precedence (i.e., in which an operation returning before another is invoked is always ordered before)
 - is the same **as if** the (read and write) operations by all processes on the same data store were executed in some **sequential order** and the operations of each individual process appear in this sequence in **its program order**

Linearizability

Linearizability (of reads and writes)

- *Linearizability*. The result of each execution is
 - equivalent to a sequential execution that respects the real-time precedence (i.e., in which an operation returning before another is invoked is always ordered before)
 - is the same as if the (read and write) operations by all processes on the same data store were executed in some sequential order and the operations of each individual process appear in this sequence in its program order
- Example 5: is this a linearizable execution? (initially $x=0$)




Linearizability

Linearizability (of reads and writes)

- *Linearizability*. The result of each execution is
 - equivalent to a sequential execution that respects the real-time precedence (i.e., in which an operation returning before another is invoked is always ordered before)
 - is the same as if the (read and write) operations by all processes on the same data store were executed in some sequential order and the operations of each individual process appear in this sequence in its program order

- Example 5: this is **not a linearizable** execution (initially $x=0$)




- $W_1(x)1$ precedes $R_0(x)0$ w.r.t real-time, thus $R_0(x)0$ cannot return the initial value of x .

Linearizability

Linearizability (of reads and writes)


- *Linearizability*. The result of each execution is
 - equivalent to a sequential execution that respects the real-time precedence (i.e., in which an operation returning before another is invoked is always ordered before)
 - is the same as if the (read and write) operations by all processes on the same data store were executed in some sequential order and the operations of each individual process appear in this sequence in its program order
- Example 6: is this a linearizable execution? (initially $x=0$)



Linearizability

Linearizability (of reads and writes)


- *Linearizability*. The result of each execution is
 - equivalent to a sequential execution that respects the real-time precedence (i.e., in which an operation returning before another is invoked is always ordered before)
 - is the same as if the (read and write) operations by all processes on the same data store were executed in some sequential order and the operations of each individual process appear in this sequence in its program order
- Example 6: this is a **linearizable** execution (initially $x=0$)



Linearizability

Linearizability (of reads and writes)


- *Linearizability*. The result of each execution is
 - equivalent to a sequential execution that respects the real-time precedence (i.e., in which an operation returning before another is invoked is always ordered before)
 - is the same as if the (read and write) operations by all processes on the same data store were executed in some sequential order and the operations of each individual process appear in this sequence in its program order
- Example 7: is this a linearizable execution? (initially $x=0$)



Linearizability

Linearizability (of reads and writes)

- *Linearizability*. The result of each execution is
 - equivalent to a sequential execution that respects the real-time precedence (i.e., in which an operation returning before another is invoked is always ordered before)
 - is the same as if the (read and write) operations by all processes on the same data store were executed in some sequential order and the operations of each individual process appear in this sequence in its program order
- Example 7: this is a **linearizable** execution (initially $x=0$)



- $R_0(x)1$ is concurrent with $W_1(x)1$ (none precedes the other w.r.t real-time), hence $R_0(x)$ can return 1 (or 0) without violating linearizability

Linearizability

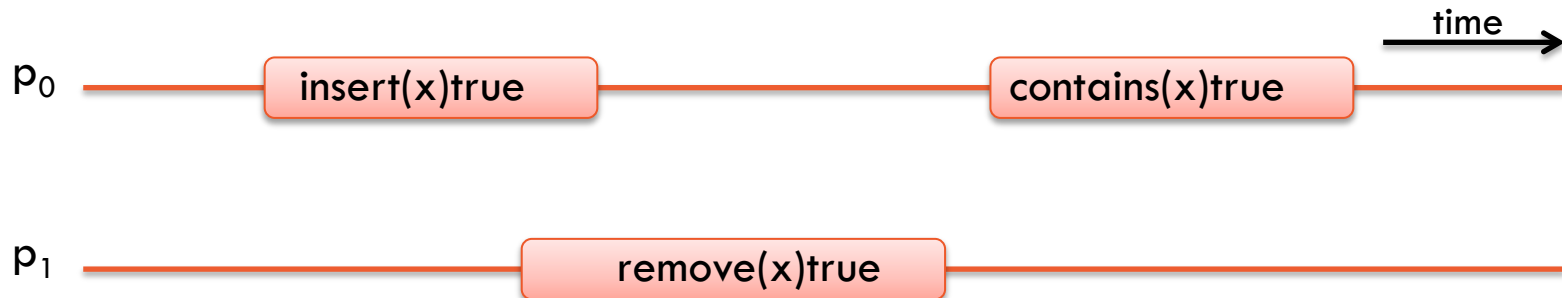
Linearizability (of higher level operations)

- Linearizability can also apply to other kind of operations.

Consider an **integer set object** that exports (as opposed to the previous register exporting read/write):

- `insert(int a)boolean b` that **adds** an element to the set if not already present and return true, returns false if already present
- `delete(int a)boolean b` that **removes** an integer from the set if it is currently present and returns true, false if not present
- `contains(int a)boolean b` that **checks** whether an integer is present in the set

- Example 8: is this a linearizable execution of set operations? (initially $x \notin \text{set}$)



Linearizability

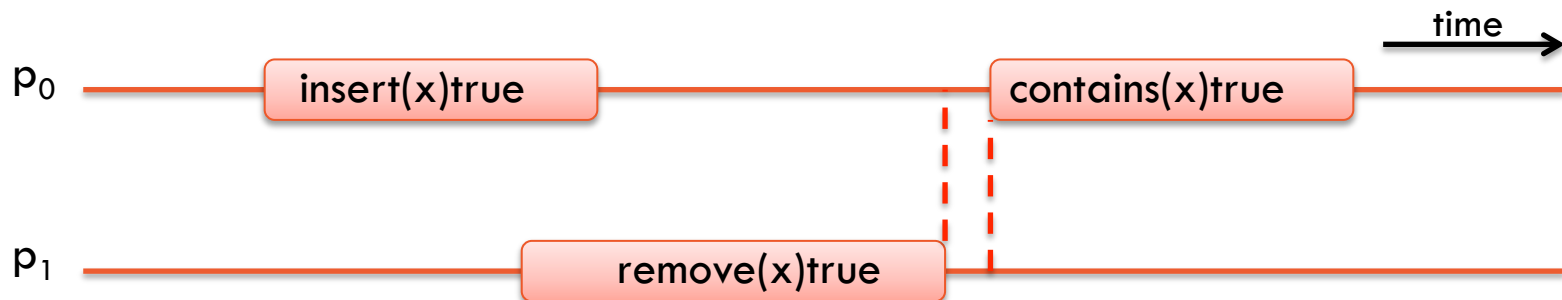
Linearizability (of higher level operations)

- Linearizability can also apply to other kind of operations.

Consider an integer set object that exports (as opposed to the previous register exporting read/write):

- `insert(int a)boolean b` that adds an element to the set if not already present and return true, returns false if already present
- `delete(int a)boolean b` that removes an integer from the set if it is currently present and returns true, false if not present
- `contains(int a)boolean b` that checks whether an integer is present in the set

- Example 8: this is **not a linearizable** execution of set operations (initially $x \notin \text{set}$)



Linearizability

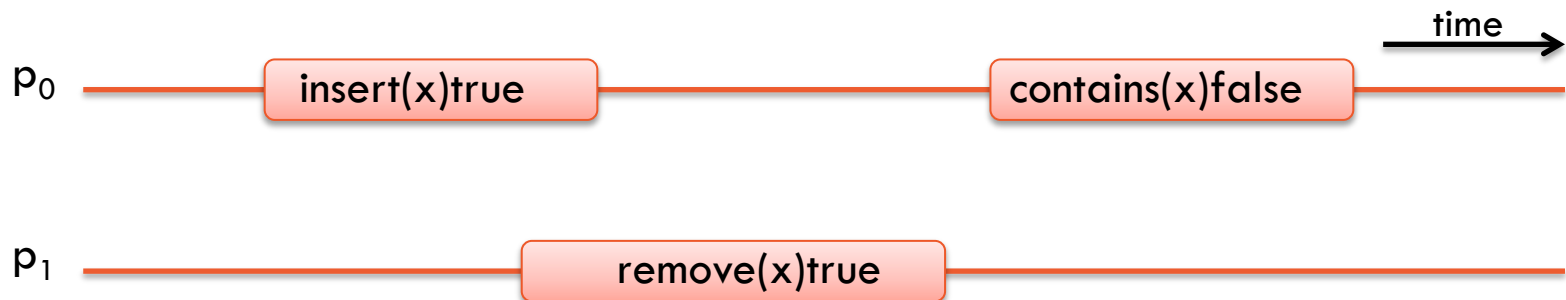
Linearizability (of higher level operations)

- Linearizability can also apply to other kind of operations.

Consider an integer set object that exports (as opposed to the previous register exporting read/write):

- `insert(int a)` boolean b that adds an element to the set if not already present and return true, returns false if already present
- `delete(int a)` boolean b that removes an integer from the set if it is currently present and returns true, false if not present
- `contains(int a)` boolean b that checks whether an integer is present in the set

- Example 9: is this a linearizable execution of set operations? (initially $x \notin \text{set}$)



Linearizability

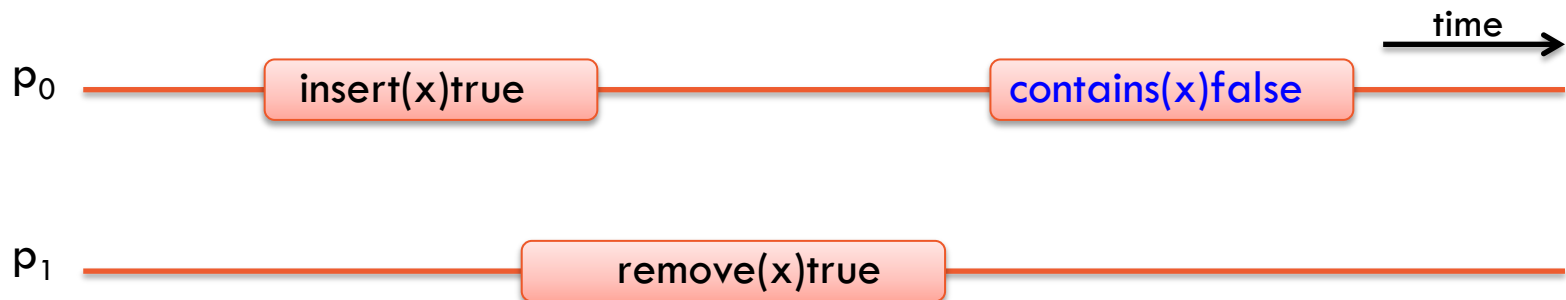
Linearizability (of higher level operations)

- Linearizability can also apply to other kind of operations.

Consider an integer set object that exports (as opposed to the previous register exporting read/write):

- `insert(int a)` boolean b that adds an element to the set if not already present and return true, returns false if already present
- `delete(int a)` boolean b that removes an integer from the set if it is currently present and returns true, false if not present
- `contains(int a)` boolean b that checks whether an integer is present in the set

- Example 9: this is a **linearizable** execution of set operations (initially $x \notin \text{set}$)



Linearizability

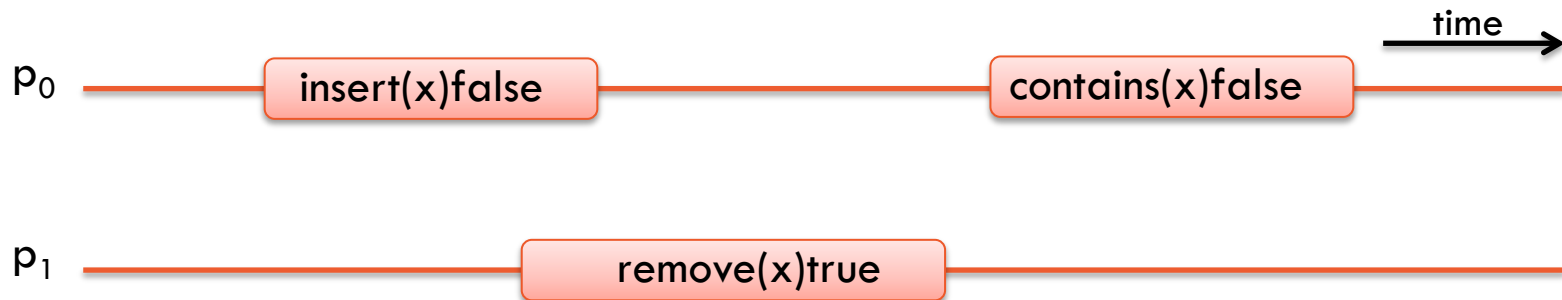
Linearizability (of higher level operations)

- Linearizability can also apply to other kind of operations.

Consider an integer set object that exports (as opposed to the previous register exporting read/write):

- `insert(int a)boolean b` that adds an element to the set if not already present and return true, returns false if already present
- `delete(int a)boolean b` that removes an integer from the set if it is currently present and returns true, false if not present
- `contains(int a)boolean b` that checks whether an integer is present in the set

- Example 10: is this a linearizable execution of set operations? (initially $x \notin \text{set}$)



Linearizability

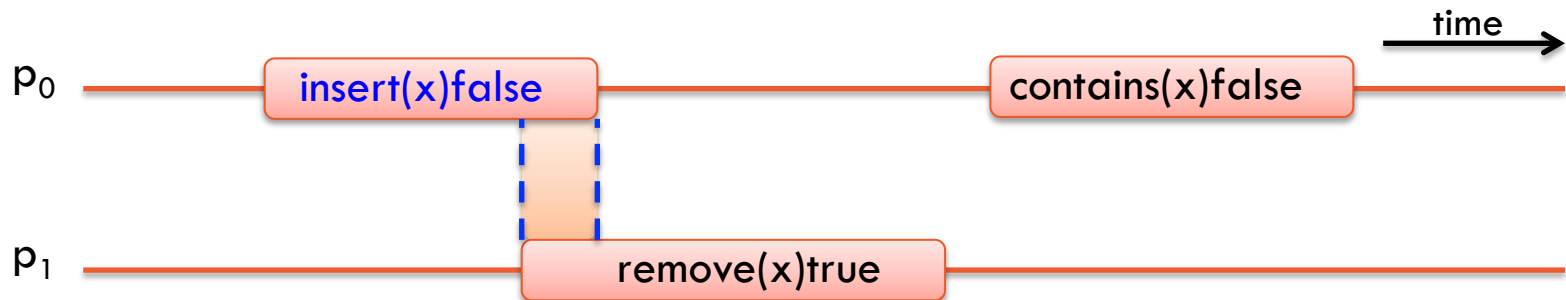
Linearizability (of higher level operations)

- Linearizability can also apply to other kind of operations.

Consider an integer set object that exports (as opposed to the previous register exporting read/write):

- `insert(int a)` boolean b that adds an element to the set if not already present and return true, returns false if already present
- `delete(int a)` boolean b that removes an integer from the set if it is currently present and returns true, false if not present
- `contains(int a)` boolean b that checks whether an integer is present in the set

- Example 10: this is a **linearizable** execution of set operations (initially $x \in \text{set}$)



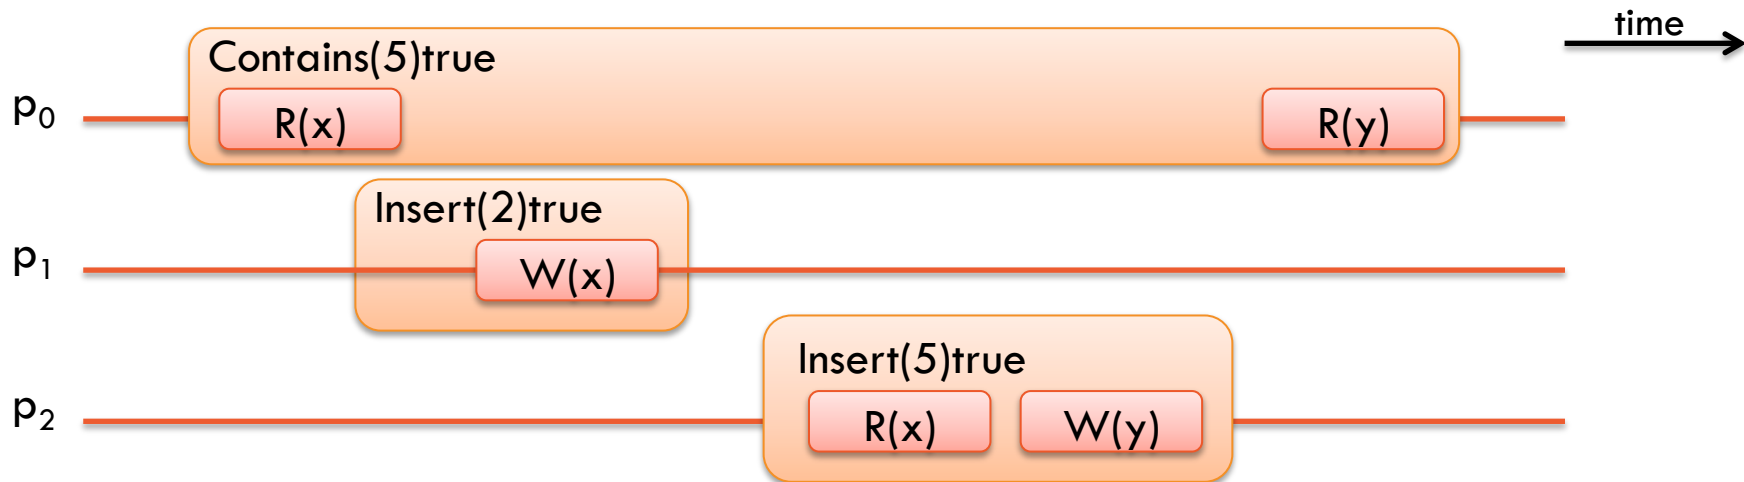
Linearizability

Linearizability (of higher level operations)

- Linearizability can also apply to other kind of operations.

Consider an integer set object that exports (as opposed to the previous register exporting read/write)

- `insert(int a)boolean b` that adds an element to the set if not already present and return true, returns false if already present
- `delete(int a)boolean b` that removes an integer from the set if it is currently present and returns true, false if not present
- `contains(int a)boolean b` that checks whether an integer is present in the set
- Example 11: this is **linearizable at the set level** but **not linearizable at the register level**



Conclusion

- There are several consistency criteria
 - Strong criteria (restrictive):
 - provides **lots of safety guarantees** (are enough for most safety requirements)
 - provides **low performance**
 - Weak criteria (not restrictive):
 - provides **less safety guarantees** (can be enough for specific applications)
 - gives **high performance**
- Some criteria strength are comparable (from the weakest to the strongest):
 - causal consistency < sequential consistency
 - **serial** \Rightarrow **sequential consistent** \Rightarrow **causally consistent**
 - **serializable** \neq **linearizability**

Backup

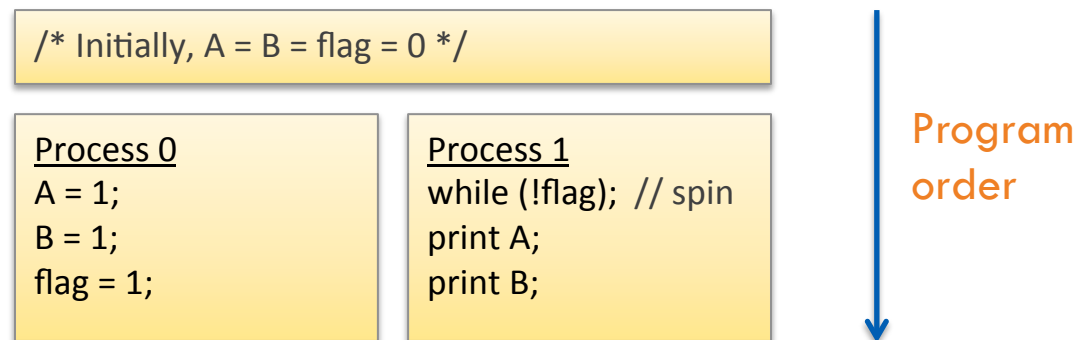


THE UNIVERSITY OF
SYDNEY

Relaxed consistency

Relaxation limitations

- Example: all orders cannot be relaxed



- Reads on A and B can be reordered and writes on A and B can be reordered
- However, writes on A and B must be before write on flag and reads of A and B must come after the read of flag.
- Solution, 2 fences are added on Compaq Alpha, IBM PowerPC, SUN RMO
 1. Between writes on A/B and write of flag and
 2. Between reads of flag and reads on A/B

Continuous Consistency



THE UNIVERSITY OF
SYDNEY

Continuous consistency

Continuous consistency

- Observation: We can actually talk about a **degree** of consistency:
 - Replicas may **differ in their numerical values**
E.g., any 2 copies of a stock market prices should not deviate by $> \$0.02$
 - Replicas may **differ in their relative staleness**
E.g., any 2 copies should not deviate by $> 0.5\%$
 - There may be differences with respect to **number and order** of performed update operations
E.g., the copy of a web cache should not be k updates far apart from the server copy
- *Consistency unit (conit)*: specifies the **data unit** over which consistency is to be measured


Continuous consistency

Continuous consistency

- Keeping track of consistency deviation (initially, all 0, grey means committed)
 - Initially committed value for x and y on both sites are 0
 - **Conit** contains the variables **x** and **y**
 - Each **replica** maintains a **vector clock**
 - B sends A operation [**<5,B>**: **x:=x+2**]; A has made this operation permanent (cannot be rolled back)
 - There is only 1 operation from B that A did not see yet:
 - numerical deviation of A is 1,
 - weight is $\max(|y_B - y_A|, |x_B - x_A|) = 5$


(where y_B is the committed value of y at B, i.e., 0 as there is no commit at B).

Replica A



Vector clock A = (15, 5)
 Order deviation = 3
 Numerical deviation = (1, 5)

Replica B

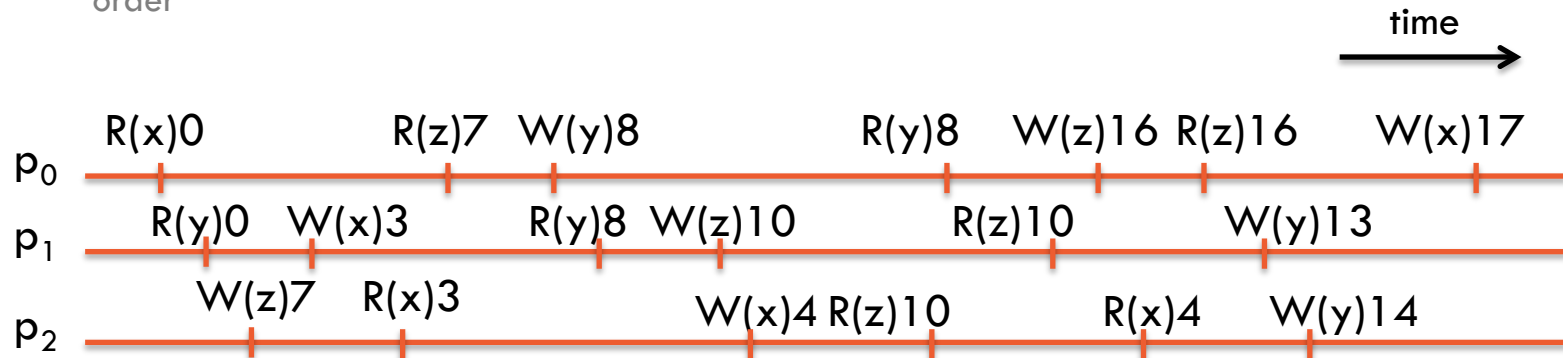


Vector clock B = (0, 11)
 Order deviation = 2
 Numerical deviation = (3, 6)

Sequential consistency

A non-sequentially consistent execution

- *Sequential consistency*. The result of each execution is
 - the same as if the (read and write) operations by all processes on the same data store were executed in some sequential order and the operations of each individual process appear in this sequence in its program order



- Example 4: this is **sequentially consistent**

- Is this equivalent to some serial execution? **No**

p₀: R(x)0; R(z)7; W(y)8; R(y)8; W(z)16; R(z)16; W(x)17 cannot be before p₁ or p₂

p₁: R(y)0; W(x)3; R(y)8; W(z)10; R(z)10; W(y)13 cannot be before p₀ (x) or p₂


p₂: W(z)7; R(x)3; W(x)4; R(z)10; R(x)4; W(y)14

p₀ writes 17 on x, but p₁ and p₂ sees different value for x; p₂ writes 14 on y but p₀, p₁ sees 8.

Sequential consistency

A non-sequentially consistent execution

- *Sequential consistency*. The result of each execution is
 - the same as if the (read and write) operations by all processes on the same data store were executed in some sequential order and the operations of each individual process appear in this sequence in its program order

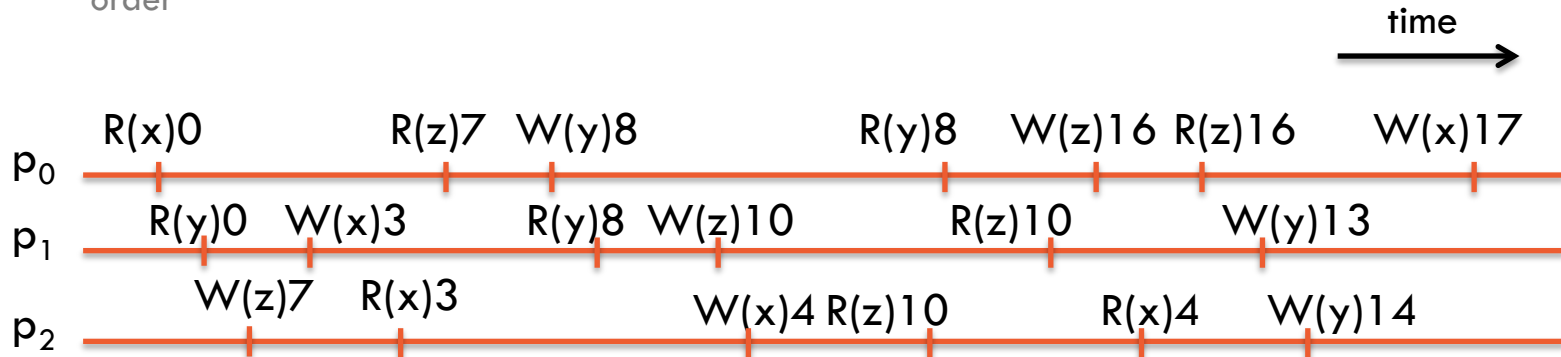


- Example 4: is the result of this execution sequentially consistent?

Sequential consistency

A non-sequentially consistent execution

- *Sequential consistency*. The result of each execution is
 - the same as if the (read and write) operations by all processes on the same data store were executed in some sequential order and the operations of each individual process appear in this sequence in its program order




- Example 4: this is **sequentially consistent**
 - Is this equivalent to some serial execution?

Sequential consistency

A non-sequentially consistent execution

- *Sequential consistency*. The result of each execution is
 - the same as if the (read and write) operations by all processes on the same data store were executed in some sequential order and the operations of each individual process appear in this sequence in its program order



- Example 4: this is **sequentially consistent**
 - Is this equivalent to some serial execution? **No**